

Parallelization in C_osmoLattice

September 2025, Daejeon

Franz R. Sattler

Bielefeld University



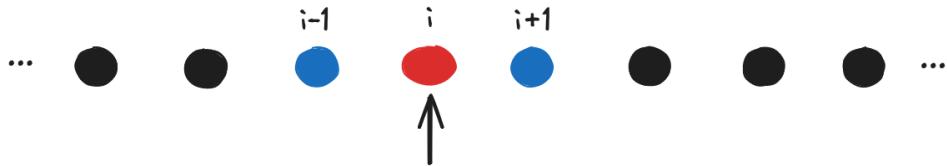
Why can we parallelize?

And why would we want to?

Example: Solving massless Klein-Gordon equation,

$$d = 1$$

$$\partial_t^2 \phi(i) = \Delta_i^+ \Delta_i^- \phi(i)$$



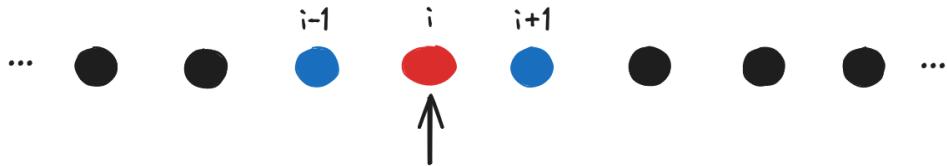
Why can we parallelize?

And why would we want to?

Example: Solving massless Klein-Gordon equation,

$$d = 1$$

$$\begin{aligned}\partial_t \pi(i) &= \Delta_i^+ \Delta_i^- \phi(i), \\ \partial_t \phi(i) &= \pi(i).\end{aligned}$$



(Leapfrog scheme.)

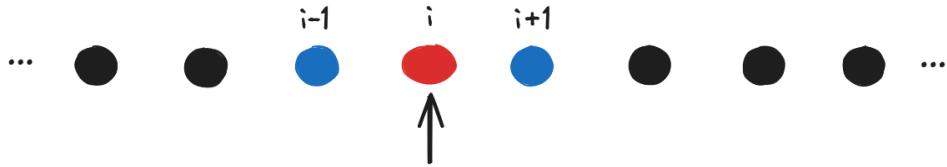
Why can we parallelize?

And why would we want to?

Example: Solving massless Klein-Gordon equation,

$$d = 1$$

$$\begin{aligned}\partial_t \pi(i) &= \Delta_i^+ \Delta_i^- \phi(i), \\ \partial_t \phi(i) &= \pi(i).\end{aligned}$$



(Leapfrog scheme.)

Stencil of update for **one** lattice site is $s = 1$

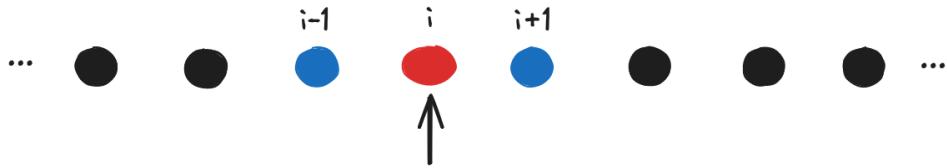
Why can we parallelize?

And why would we want to?

Example: Solving massless Klein-Gordon equation,

$$d = 1$$

$$\begin{aligned}\partial_t \pi(i) &= \Delta_i^+ \Delta_i^- \phi(i), \\ \partial_t \phi(i) &= \pi(i).\end{aligned}$$



(Leapfrog scheme.)

Stencil of update for **one** lattice site is $s = 1$

→ Only local information needed (neighbours).

Why can we parallelize?

And why would we want to?

Example: Solving massless Klein-Gordon equation,

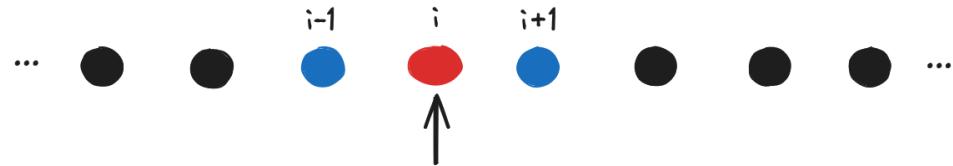
$$d = 1$$

$$\begin{aligned}\partial_t \pi(i) &= \Delta_i^+ \Delta_i^- \phi(i), \\ \partial_t \phi(i) &= \pi(i).\end{aligned}$$

(Leapfrog scheme.)

Stencil of update for **one** lattice site is $s = 1$

→ Only local information needed (neighbours).



We could compute all lattice sites independently!

See lecture on Friday!

Why can we parallelize?

And why would we want to?

Example: Solving massless Klein-Gordon equation,

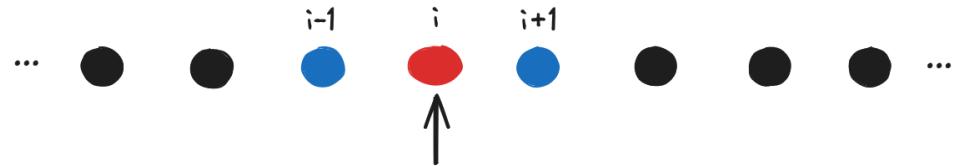
$$d = 1$$

$$\begin{aligned}\partial_t \pi(i) &= \Delta_i^+ \Delta_i^- \phi(i), \\ \partial_t \phi(i) &= \pi(i).\end{aligned}$$

(Leapfrog scheme.)

Stencil of update for **one** lattice site is $s = 1$

→ Only local information needed (neighbours).



We could compute all lattice sites independently!

See lecture on Friday!

Why can we parallelize?

And why would we want to?

Example: Solving massless Klein-Gordon equation,

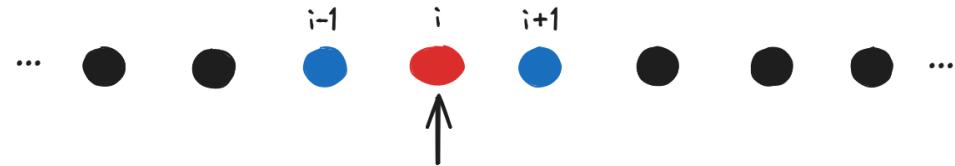
$$d = 1$$

$$\begin{aligned}\partial_t \pi(i) &= \Delta_i^+ \Delta_i^- \phi(i), \\ \partial_t \phi(i) &= \pi(i).\end{aligned}$$

(Leapfrog scheme.)

Stencil of update for **one** lattice site is $s = 1$

→ Only local information needed (neighbours).



We could compute all lattice sites independently!

See lecture on Friday!

Less granular: split **sub-regions** of lattice across many computers (**nodes**)

Parallelization

of CosmoLattice

simulations requires to split both computation and data across cores.

Parallelization

of CosmoLattice

simulations requires to split both **computation** and **data** across **cores**.

Cores: **Nodes** (distributed) and **Threads** (shared).

Parallelization

of CosmoLattice

simulations requires to split both **computation** and **data** across **cores**.

Cores: **Nodes** (distributed) and **Threads** (shared).

Type	distributed	shared
------	--------------------	---------------

Data	split between nodes	shared by all threads
------	----------------------------	------------------------------

Computation	split between nodes	split between threads
-------------	----------------------------	------------------------------

Parallelization

of CosmoLattice

simulations requires to split both computation and data across cores.

$$\text{speedup} = \frac{1}{(1 - \alpha) + \frac{\alpha}{n_{\text{cores}}}}$$

α = part of the code which runs in parallel

n_{cores} = speedup of the parallel part

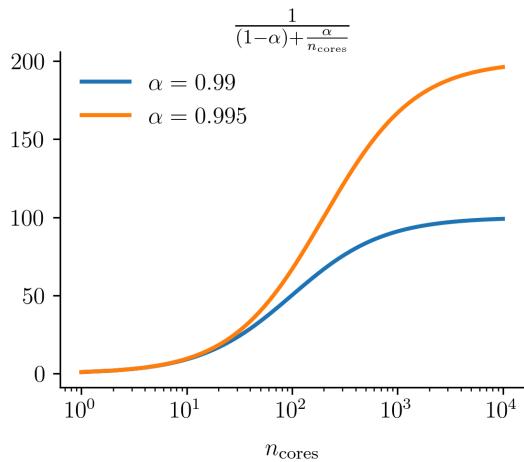
CosmoLattice: $\alpha \gtrsim 0.99$

(see manual)

Parallelization

of CosmoLattice

simulations requires to split both computation and data across cores.



$$\text{speedup} = \frac{1}{(1 - \alpha) + \frac{\alpha}{n_{\text{cores}}}}$$

CosmoLattice: $\alpha \gtrsim 0.99$
(see manual)

α = part of the code which runs in parallel
 n_{cores} = speedup of the parallel part

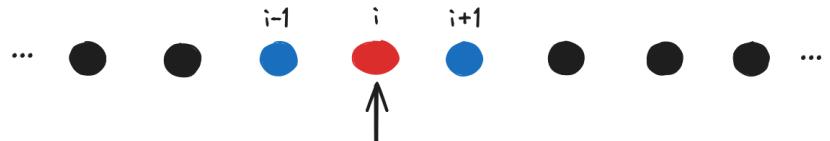
Why can we parallelize?

And why would we want to?

Example: Solving massless Klein-Gordon equation,

$$d = 1$$

$$\begin{aligned}\partial_t \pi(i) &= \Delta_i^+ \Delta_i^- \phi(i), \\ \partial_t \phi(i) &= \pi(i).\end{aligned}$$



(Leapfrog scheme.)

Stencil of update for **one** lattice site is $s = 1$

- Only local information needed (neighbours).
- Each site can be updated independently.

Why can we parallelize?

And why would we want to?

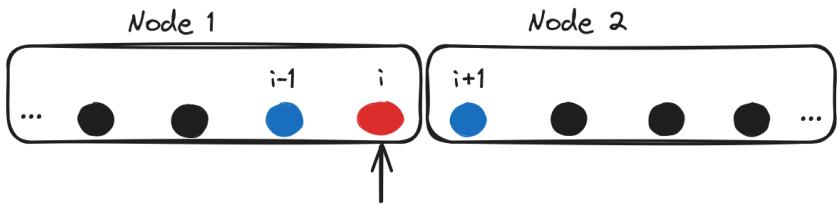
Example: Solving massless Klein-Gordon equation,
 $d = 1$

$$\begin{aligned}\partial_t \pi(i) &= \Delta_i^+ \Delta_i^- \phi(i), \\ \partial_t \phi(i) &= \pi(i).\end{aligned}$$

(Leapfrog scheme.)

Stencil of update for **one** lattice site is $s = 1$

- Only local information needed (neighbours).
- Each site can be updated independently.



Why can we parallelize?

And why would we want to?

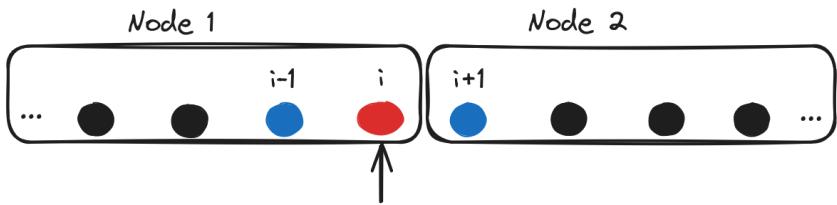
Example: Solving massless Klein-Gordon equation,
 $d = 1$

$$\begin{aligned}\partial_t \pi(i) &= \Delta_i^+ \Delta_i^- \phi(i), \\ \partial_t \phi(i) &= \pi(i).\end{aligned}$$

(Leapfrog scheme.)

Stencil of update for **one** lattice site is $s = 1$

- Only local information needed (neighbours).
- Each site can be updated independently.



Problem: Data is missing on node 1

Why can we parallelize?

And why would we want to?

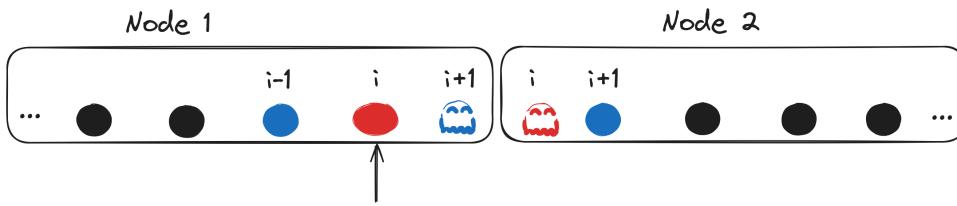
Example: Solving massless Klein-Gordon equation,
 $d = 1$

$$\begin{aligned}\partial_t \pi(i) &= \Delta_i^+ \Delta_i^- \phi(i), \\ \partial_t \phi(i) &= \pi(i).\end{aligned}$$

(Leapfrog scheme.)

Stencil of update for **one** lattice site is $s = 1$

- Only local information needed (neighbours).
- Each site can be updated independently.



Solution: Use ghosts.

Ghosts are local copies of data on other nodes.

Why can we parallelize?

And why would we want to?

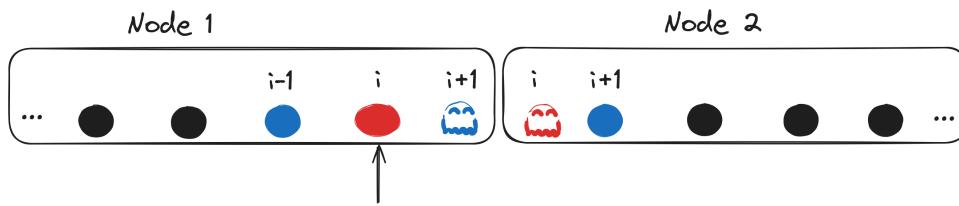
Example: Solving massless Klein-Gordon equation,
 $d = 1$

$$\begin{aligned}\partial_t \pi(i) &= \Delta_i^+ \Delta_i^- \phi(i), \\ \partial_t \phi(i) &= \pi(i).\end{aligned}$$

(Leapfrog scheme.)

Stencil of update for **one** lattice site is $s = 1$

- Only local information needed (neighbours).
- Each site can be updated independently.



Solution: Use ghosts.

Ghosts are local copies of data on other nodes.

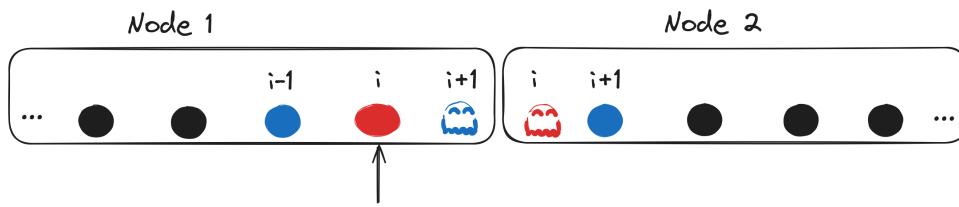
Need to update ghosts after every time-step.

Data communication

The standard for communication in distributed-memory applications:

Message Passing Interface (MPI)

Exchange ghost data between **nodes** over the **network** automatically if anything changes.

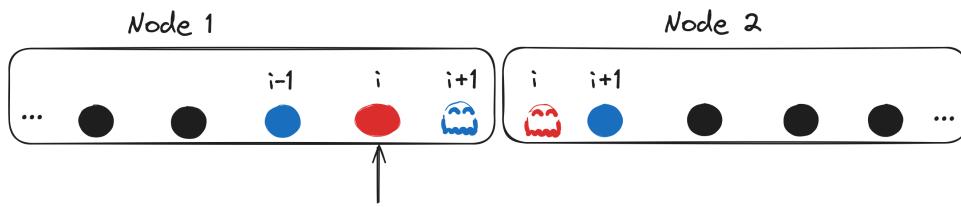


Data communication

The standard for communication in distributed-memory applications:

Message Passing Interface (MPI)

Exchange ghost data between **nodes** over the **network** automatically if anything changes.



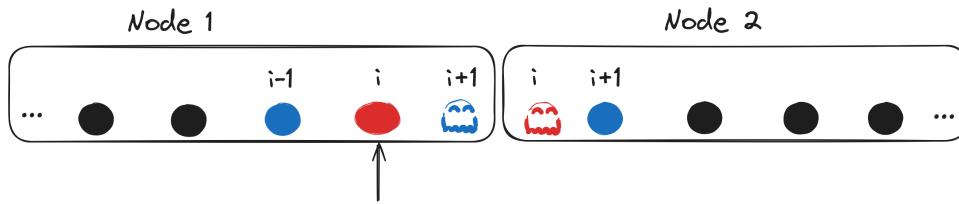
CosmoLattice does this automatically under the hood!

Data communication

The standard for communication in distributed-memory applications:

Message Passing Interface (MPI)

Exchange ghost data between **nodes** over the **network** automatically if anything changes.



CosmoLattice does this automatically under the hood!

How to turn it on?

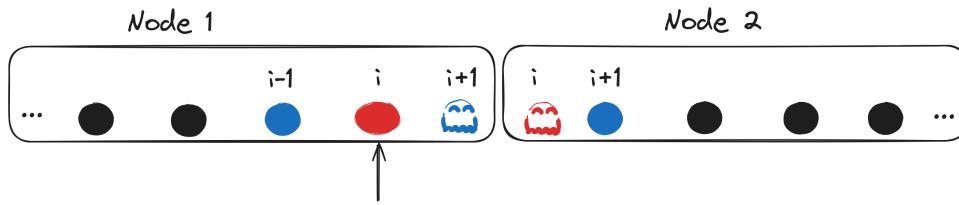
```
$ cd build/  
$ cmake -DMPI=ON ..  
$ make  
$ mpirun -n 16 ./lphi4 input=./lphi4.in
```

Data communication

The standard for communication in distributed-memory applications:

Message Passing Interface (MPI)

Exchange ghost data between **nodes** over the **network** automatically if anything changes.



CosmoLattice does this automatically under the hood!

How to turn it on?

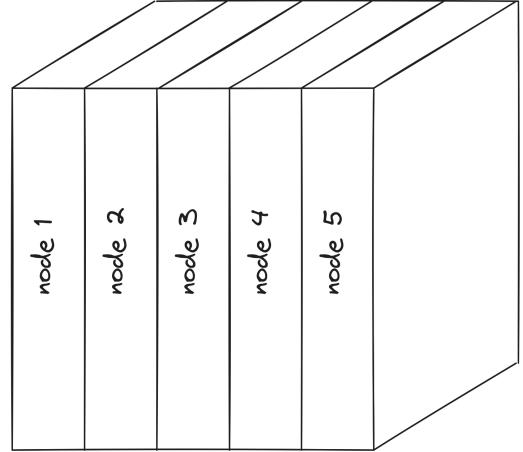
```
$ cd build/  
$ cmake -DMPI=ON ..  
$ make  
$ mpirun -n 16 ./lphi4 input=./lphi4.in
```

You may need to install fftw3 with MPI support.

```
$ sudo apt-get install libfftw3-mpi-dev
```

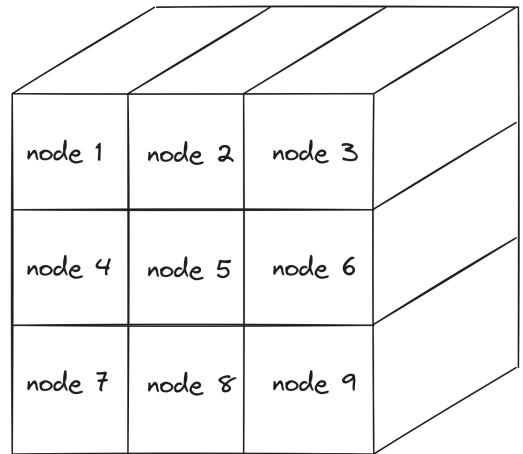
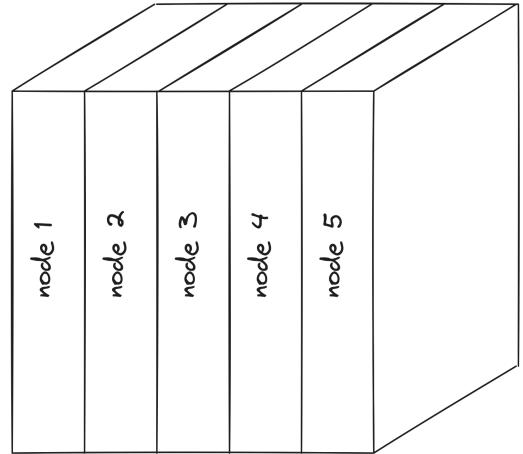
FFT parallelization

- FFTW supports parallelization along 1 direction.



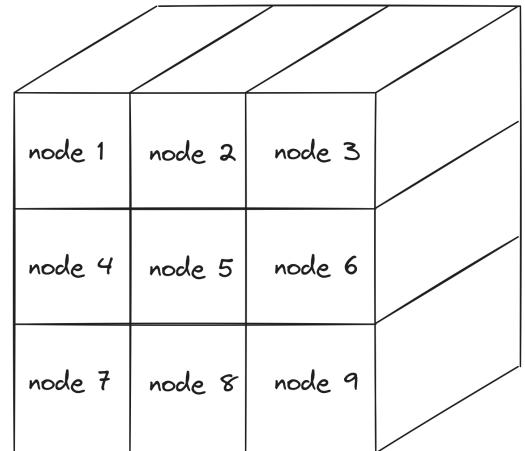
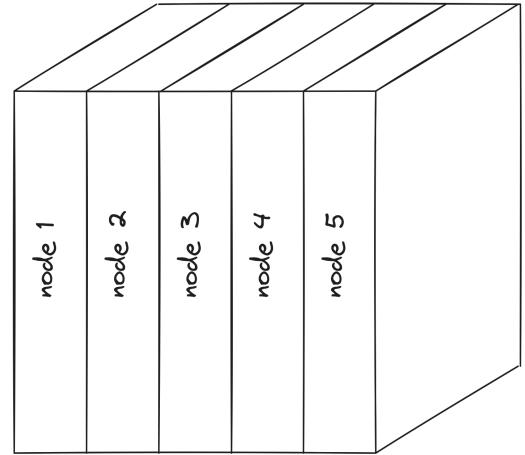
FFT parallelization

- FFTW supports parallelization along 1 direction.
- Improved scaling for many nodes: PFFT allows for parallelization in any dimensions.



FFT parallelization

- FFTW supports parallelization along 1 direction.
- Improved scaling for many nodes: PFFT allows for parallelization in any dimensions.
- To keep data transfer due to ghost exchange manageable, parallelization along $d - 1$ directions.



FFT parallelization

- FFTW supports parallelization along 1 direction.
- Improved scaling for many nodes: PFFT allows for parallelization in any dimensions.
- To keep data transfer due to ghost exchange manageable, parallelization along $d - 1$ directions.

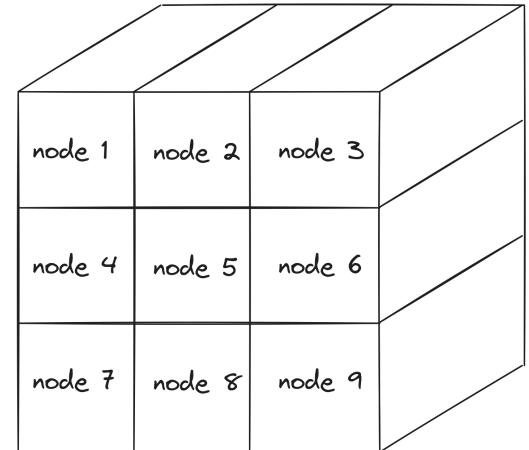
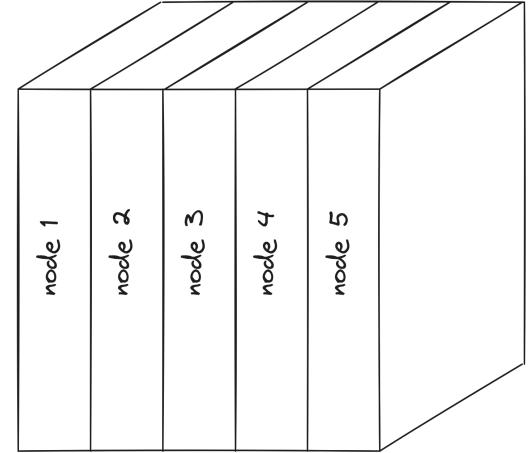
1D

$$N = n_p * m$$

$$N = 50$$

2D

$$\begin{aligned} N &= n_p^{(1)} * m \\ &= n_p^{(2)} * m \end{aligned}$$



FFT parallelization

- FFTW supports parallelization along 1 direction.
- Improved scaling for many nodes: PFFT allows for parallelization in any dimensions.
- To keep data transfer due to ghost exchange manageable, parallelization along $d - 1$ directions.

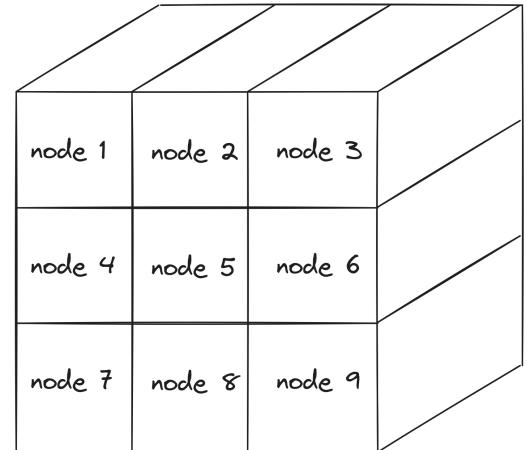
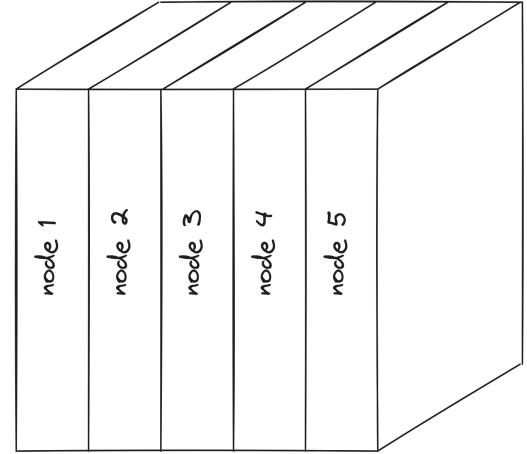
1D

$$N = n_p * m$$

$$N = 50$$

2D

$$\begin{aligned} N &= n_p^{(1)} * m \\ &= n_p^{(2)} * m \end{aligned}$$



FFT parallelization

- FFTW supports parallelization along 1 direction.
- Improved scaling for many nodes: PFFT allows for parallelization in any dimensions.
- To keep data transfer due to ghost exchange manageable, parallelization along $d - 1$ directions.

1D

$$N = n_p * m$$

25 nodes.

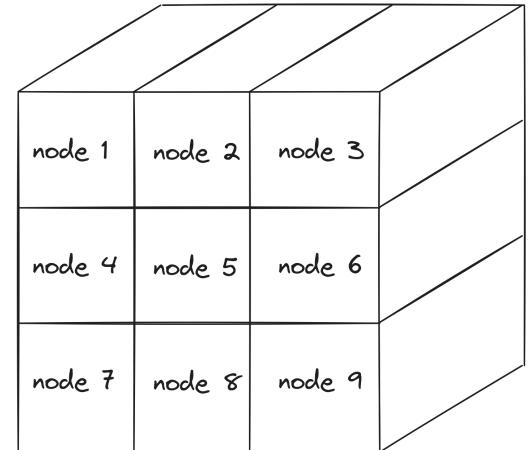
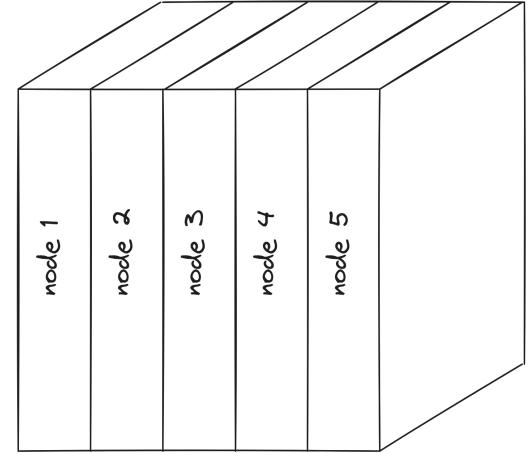
$$N = 50$$

Maximum parallelization

2D

$$\begin{aligned} N &= n_p^{(1)} * m \\ &= n_p^{(2)} * m \end{aligned}$$

$25^2 = 625$ nodes.



FFT parallelization

- FFTW supports parallelization along 1 direction.
- Improved scaling for many nodes: PFFT allows for parallelization in any dimensions.
- To keep data transfer due to ghost exchange manageable, parallelization along $d - 1$ directions.

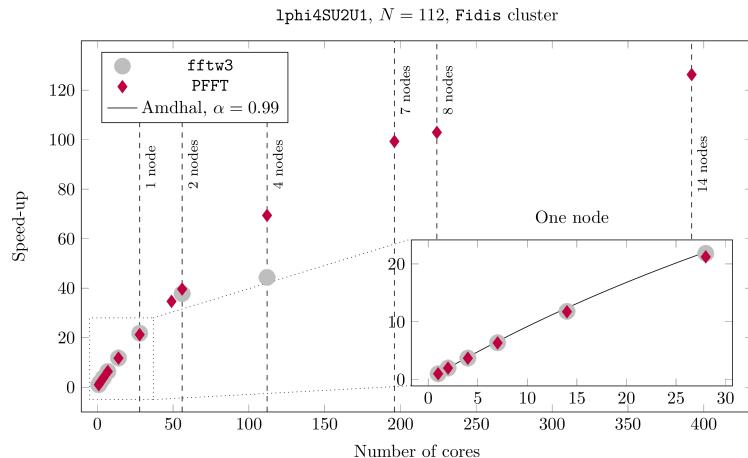
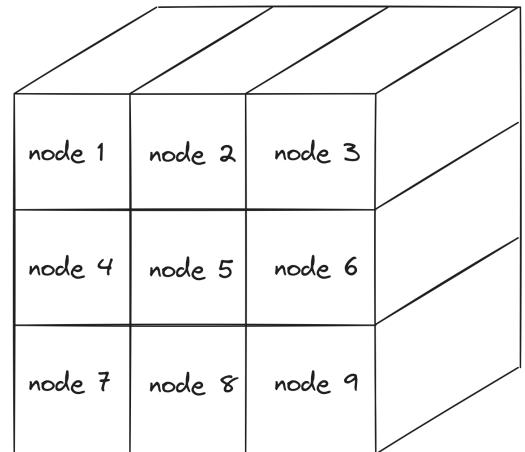
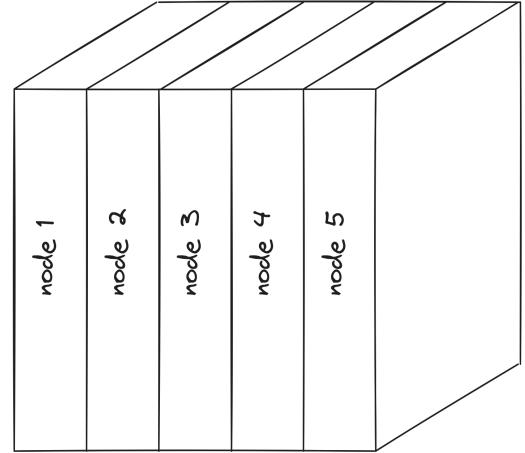
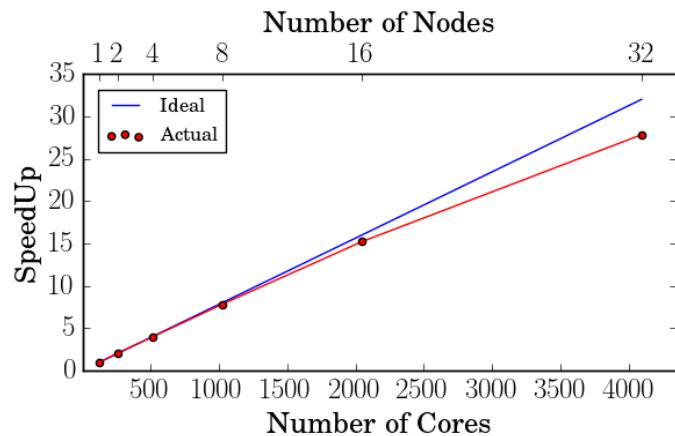
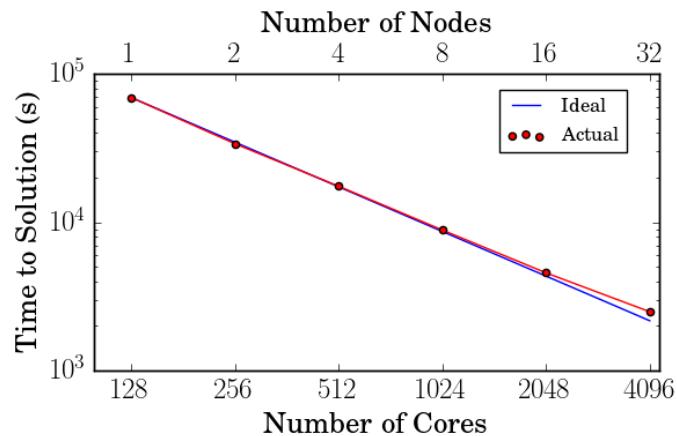


Figure 3: Speed up factor in parallelized simulations as a number of cores (tested on the Gacrux cluster from the EPFL HPC center SCITAS, Switzerland).



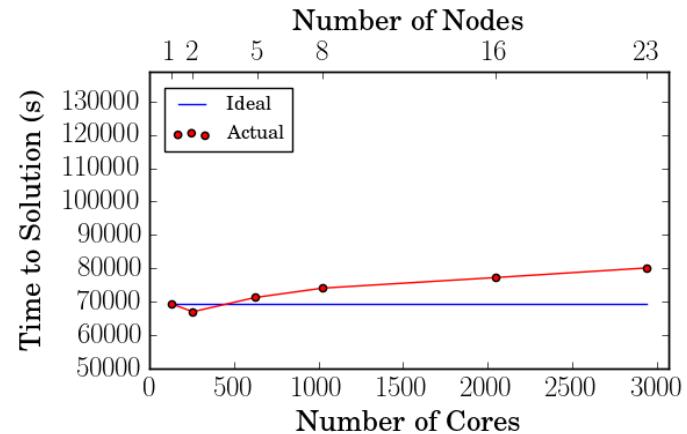
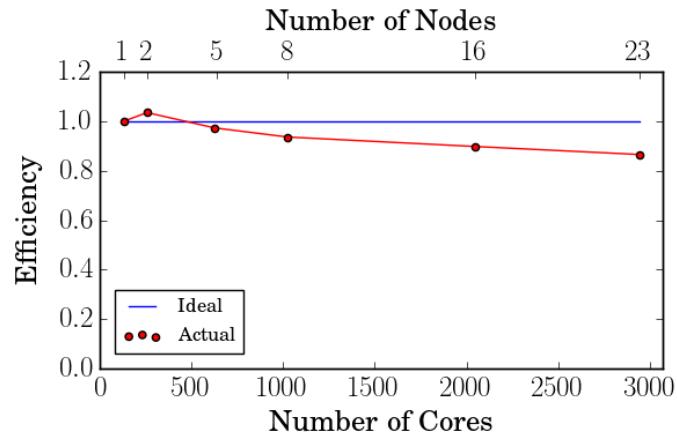
Going to larger clusters

Strong scaling
(same lattice size, more cores)



Going to larger clusters

Weak scaling
(lattice size ~ cores)



Questions?

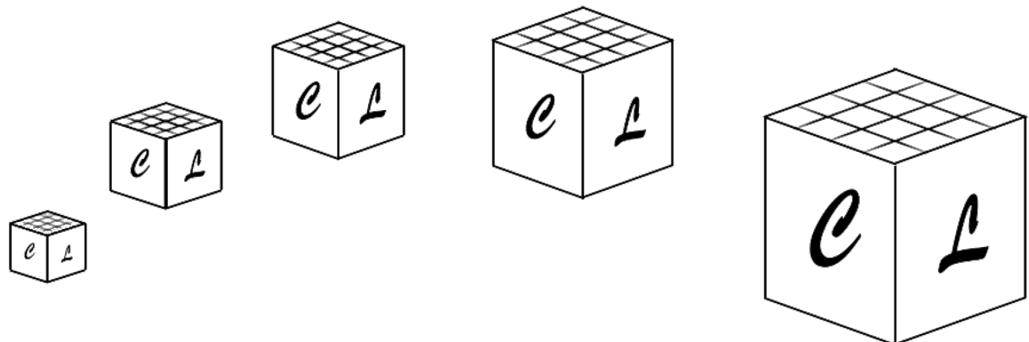
Tomorrow: Shared-memory parallelization with GPUs.

CosmoLattice on GPUs

September 2025, Daejeon

Franz R. Sattler

Bielefeld University



Current computers

can be broadly said to have two main processing units:

Current computers

can be broadly said to have two main

processing units: **CPU**s (central processing units)

Current computers

can be broadly said to have two main

processing units: **CPU**s (central processing units) and **GPU**s (graphical processing units)

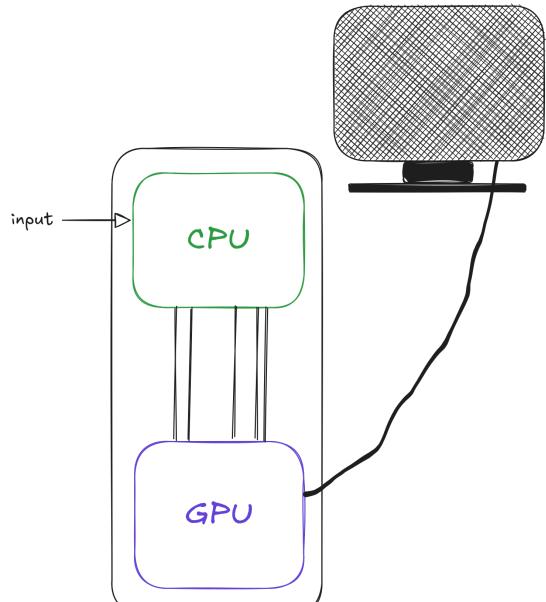
Current computers

can be broadly said to have two main

processing units: **CPU**s (central processing units) and **GPU**s (graphical processing units)

CPUs: for **OS, computation, general applications**.

GPUs: dedicated just for **video and graphics** applications.



A consumer machine

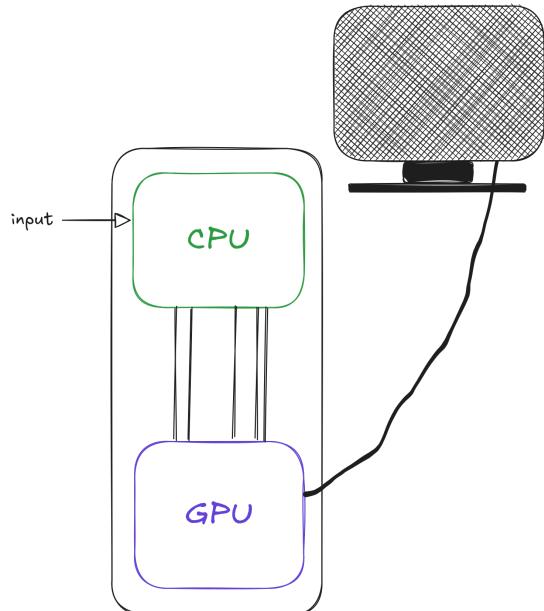
Current computers

can be broadly said to have two main processing units: **CPU**s (central processing units) and **GPU**s (graphical processing units)

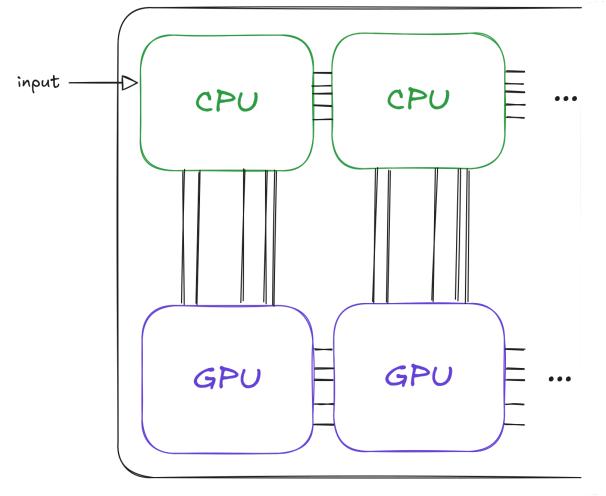
CPUs: for **OS, computation, general applications.**

GPUs: dedicated just for **video and graphics** applications.

Current (heterogeneous) clusters have both **CPU**s and **GPU**s for **computations**.



A consumer machine



A typical heterogeneous computing cluster

Why use GPUs for lattice?

Lattice points independently computed & updated → Limit of threads is number of lattice sites!

Why use GPUs for lattice?

Lattice points independently computed & updated → Limit of threads is number of lattice sites!

	CPU	GPU	
■ AMD EPYC 7763: 64			■ Nvidia H100: ~15000
■ Intel Xeon 6148 (Skylake): 20	Cores/Node	$\mathcal{O}(10 - 100)$	■ Nvidia 4070 RTX mobile: ~5000
■ AMD Ryzen 9 7945HX: 16	Clock speed	~ 3 GHz	~ 1.5 GHz

Why use GPUs for lattice?

Lattice points independently computed & updated → Limit of threads is number of lattice sites!

	CPU	GPU	
■ AMD EPYC 7763: 64	Cores/Node	$\mathcal{O}(10 - 100)$	$\mathcal{O}(10000)$
■ Intel Xeon 6148 (Skylake): 20	Clock speed	~ 3 GHz	~ 1.5 GHz

■ Nvidia H100: **~15000**
■ Nvidia 4070 RTX mobile: **~5000**

CPU: Low parallelization, high clock speed

GPU: High parallelization, moderate clock speed

→ CosmoLattice on GPUs has the potential for *massive parallelism* with $\gg 10^5$ simultaneous operations.

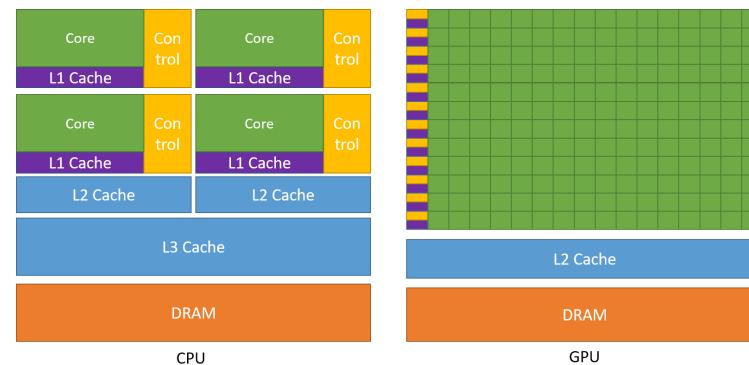
Why use GPUs for lattice?

Lattice points independently computed & updated → Limit of threads is number of lattice sites!

	CPU	GPU
Cache/Thread	64KB / 16MB	1KB
Local Cache	64MB	256KB / 50MB

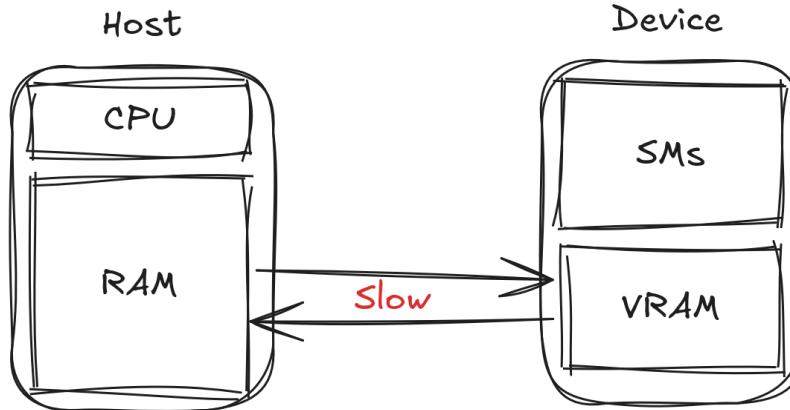
CPU: Thread-constrained

GPU: Memory-constrained



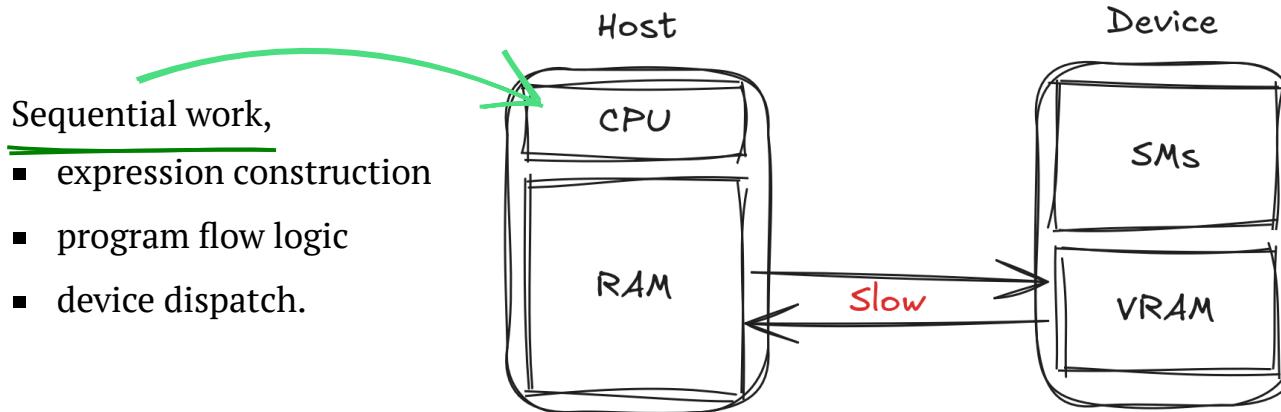
Redesigning TempLat for GPUs

Device-centric programming.



Redesigning TempLat for GPUs

Device-centric programming.

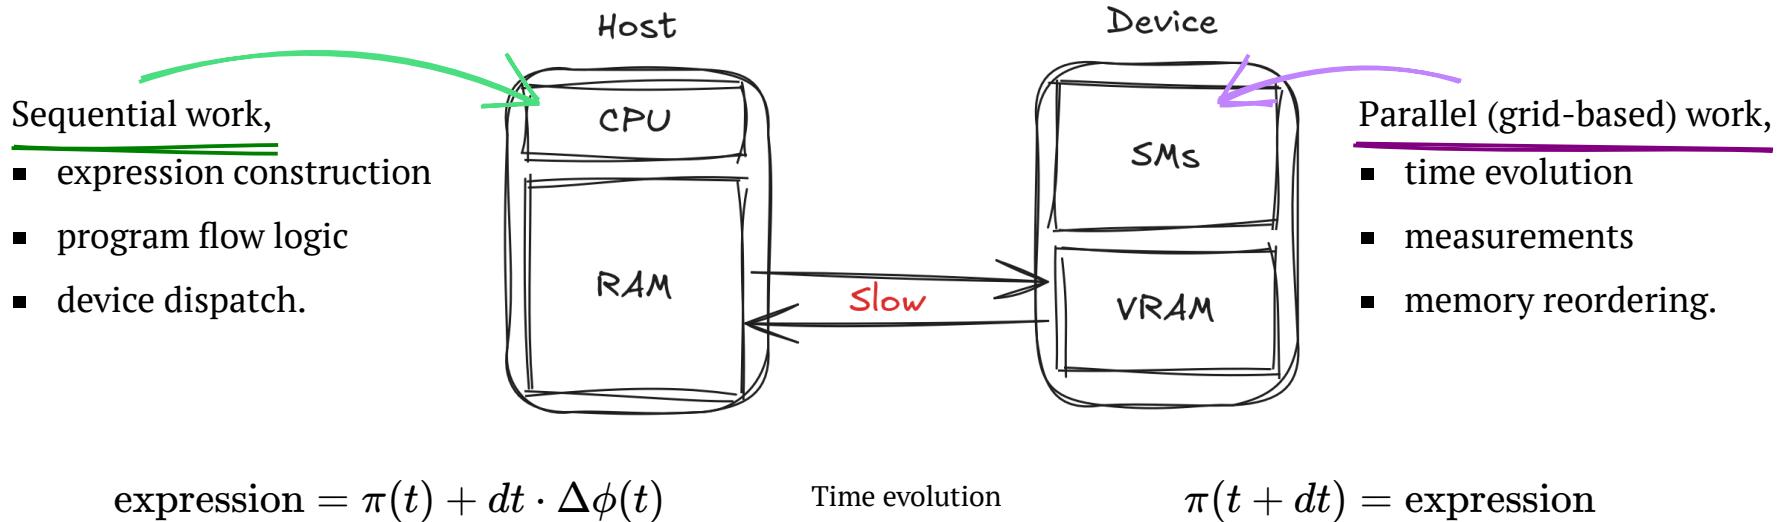


$$\text{expression} = \pi(t) + dt \cdot \Delta\phi(t)$$

Time evolution

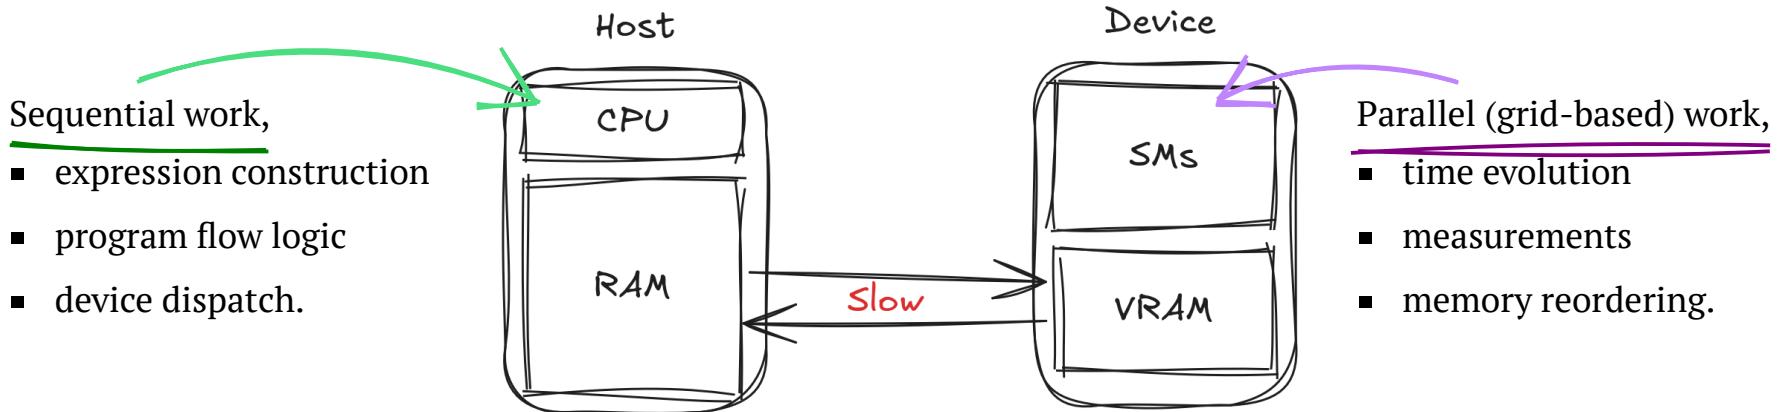
Redesigning TempLat for GPUs

Device-centric programming.



Redesigning TempLat for GPUs

Device-centric programming.



$$\text{expression} = \pi(t) + dt \cdot \Delta\phi(t)$$

Time evolution

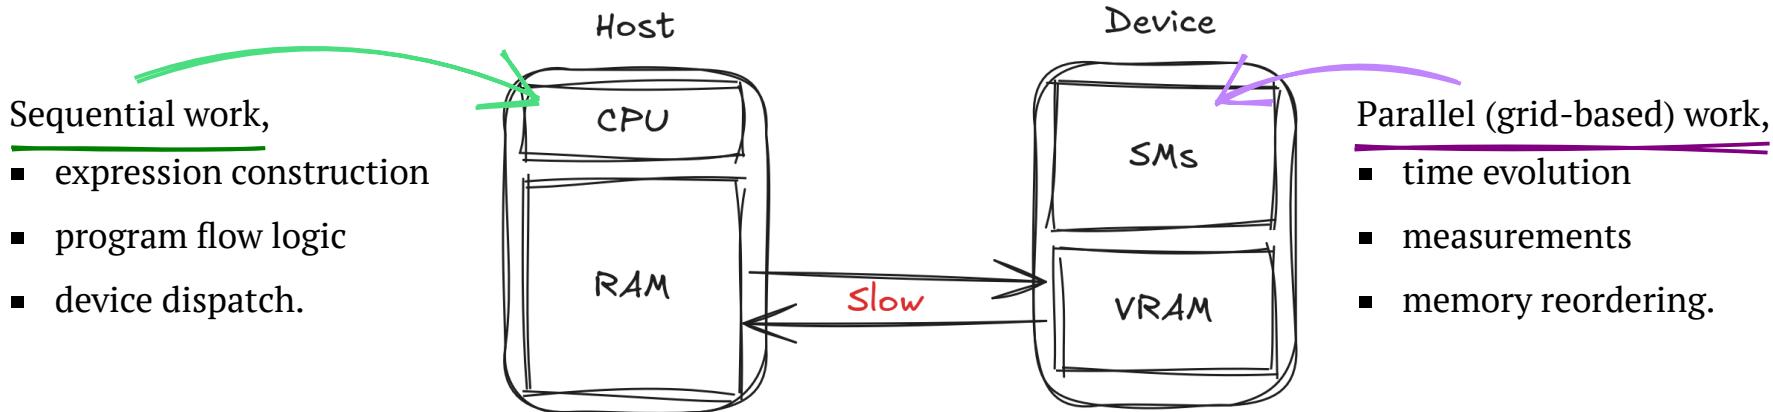
$$\pi(t + dt) = \text{expression}$$

```
auto functor = DEVICE_LAMBDA(device::IdxArray<NDim> idx,
                             double& max) {
    // ...
};
double maximum = 0.;
```

Maximum

Redesigning TempLat for GPUs

Device-centric programming.



$$\text{expression} = \pi(t) + dt \cdot \Delta\phi(t)$$

Time evolution

$$\pi(t + dt) = \text{expression}$$

```
auto functor = DEVICE_LAMBDA(device::IdxArray<NDim> idx,
                             double& max) {
    // ...
};
double maximum = 0.;
```

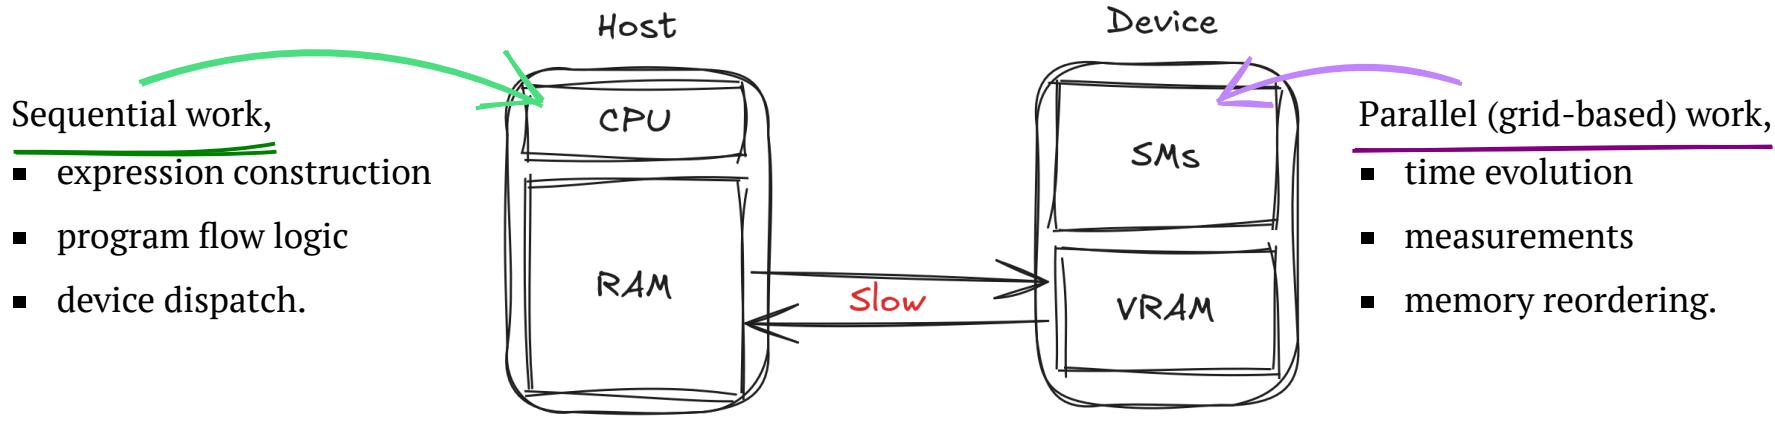
Maximum

```
device::iteration::reduce("Maximum", functor, maximum);
```



Redesigning TempLat for GPUs

Device-centric programming.



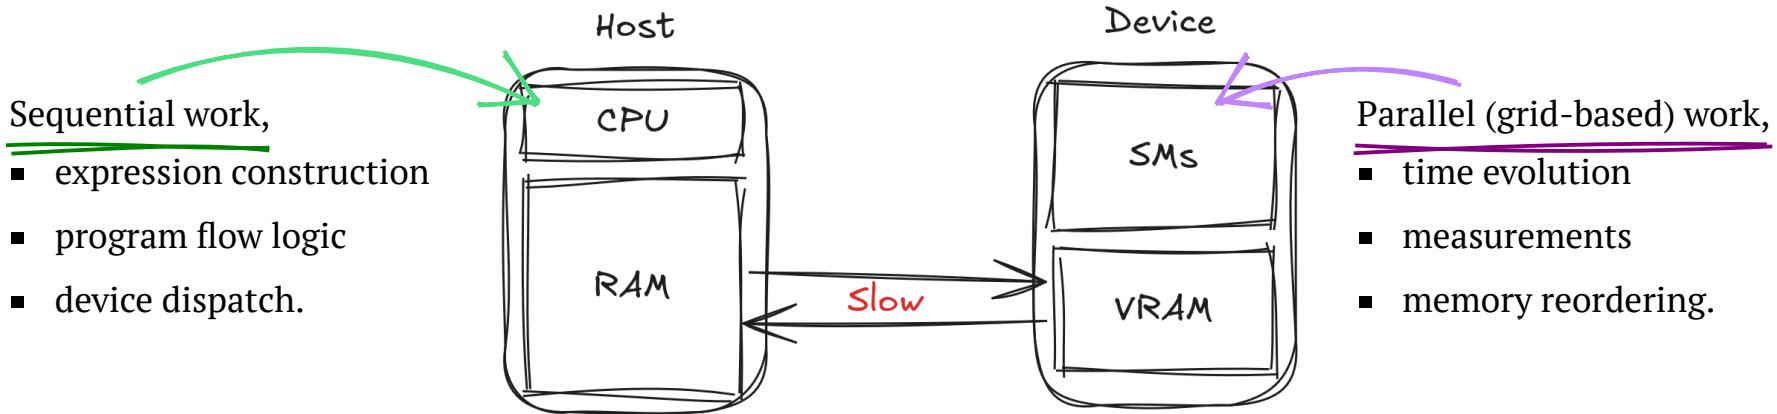
Standard C++ on CPU

Hardware-dependent

- Nvidia: CUDA
- AMD: ROCM
- Intel: SYCL
- shared-memory CPUs
- FPGAs

Redesigning TempLat for GPUs

Device-centric programming.



Standard C++ on CPU

Backends

Abstracted away in TempLat

- Kokkos
- Sequential STL (2020/2023)
- ...

- `device::iterate::foreach`
- `device::iterate::reduce`
- `device::memory::copyHostToDevice`
- ...

Does this make CosmoLattice harder to use?

Does this make CosmoLattice harder to use?

No, but...

Does this make CosmoLattice harder to use?

No, but...

Model file

```
1 public:  
2  
3     MODELNAME(ParameterParser &parser, RunParameters<double> &runPar,  
4         std::shared_ptr<MemoryToolBox> toolBox)  
5     ...
```

Does this make CosmoLattice harder to use?

No, but...

Model file

```
1 public:  
2     static constexpr size_t NDim = Model<MODELNAME>::NDim;  
3  
4     MODELNAME(ParameterParser &parser, RunParameters<double> &runPar,  
5         std::shared_ptr<MemoryToolBox<NDim>> toolBox)  
6     ...
```

Does this make CosmoLattice harder to use?

No, but...

Model file

```
1 public:  
2     static constexpr size_t NDim = Model<MODELNAME>::NDim;  
3  
4     MODELNAME(ParameterParser &parser, RunParameters<double> &runPar,  
5         std::shared_ptr<MemoryToolBox<NDim>> toolBox)  
6     ...  
7 
```

TempLat

```
1 vType computeConfigurationSpace() {  
2     vType localResult{};  
3  
4     auto& it = mT.getToolBox()->itX();  
5     for(it.begin();it.end();++it)  
6     {  
7         const ptrdiff_t i = it();  
8         localResult += GetValue::get(mT,i);  
9     }  
10    return mWorkspace;  
11 }  
12 }
```

Does this make CosmoLattice harder to use?

No, but...

Model file

```
1 public:  
2     static constexpr size_t NDim = Model<MODELNAME>::NDim;  
3  
4     MODELNAME(ParameterParser &parser, RunParameters<double> &runPar,  
5             std::shared_ptr<MemoryToolBox<NDim>> toolBox)  
6     ...  
7 
```

TempLat

```
1 vType computeConfigurationSpace() {  
2     auto functor = DEVICE_CLASS_LAMBDA(const device::IdxArray<NDim> &idx,  
3                                         vType &update) {  
4         device::apply([&](const auto &...args) {  
5             update += GetValue::get(mT, args...);  
6         },  
7             idx);  
8     };  
9  
10    vType localResult{};  
11    device::iteration::reduce("Averager", cLayout, functor, localResult);  
12    return localResult;  
13 }
```

Does this make CosmoLattice harder to use?

No, but...

Model file

```
1 public:  
2     static constexpr size_t NDim = Model<MODELNAME>::NDim;  
3  
4     MODELNAME(ParameterParser &parser, RunParameters<double> &runPar,  
5             std::shared_ptr<MemoryToolBox<NDim>> toolBox)  
6     ...  
7 
```

Building

Cosmolattice up to now:

```
1 $ cmake .. -DMODEL=lphi4  
2 ...  
3 
```

TempLat

```
1 vType computeConfigurationSpace() {  
2     auto functor = DEVICE_CLASS_LAMBDA(const device::IdxArray<NDim> &idx,  
3                                         vType &update) {  
4         device::apply([&](const auto &...args) {  
5             update += GetValue::get(mT, args...);  
6         },  
7             idx);  
8     };  
9  
10    vType localResult{};  
11    device::iteration::reduce("Averager", cLayout, functor, localResult);  
12    return localResult;  
13 }
```

Does this make CosmoLattice harder to use?

No, but...

Model file

```
1 public:  
2     static constexpr size_t NDim = Model<MODELNAME>::NDim;  
3  
4     MODELNAME(ParameterParser &parser, RunParameters<double> &runPar,  
5             std::shared_ptr<MemoryToolBox<NDim>> toolBox)  
6     ...
```

Building

New version: CUDA is detected automatically:

```
1 $ cmake .. -DMODEL=lphi4  
2 ...
```

TempLat

```
1 vType computeConfigurationSpace() {  
2     auto functor = DEVICE_CLASS_LAMBDA(const device::IdxArray<NDim> &idx,  
3                                         vType &update) {  
4         device::apply([&](const auto &...args) {  
5             update += GetValue::get(mT, args...);  
6         },  
7                     idx);  
8     };  
9  
10    vType localResult{};  
11    device::iteration::reduce("Averager", cLayout, functor, localResult);  
12    return localResult;  
13 }
```

Does this make CosmoLattice harder to use?

No, but...

Model file

```
1 public:  
2     static constexpr size_t NDim = Model<MODELNAME>::NDim;  
3  
4     MODELNAME(ParameterParser &parser, RunParameters<double> &runPar,  
5             std::shared_ptr<MemoryToolBox<NDim>> toolBox)  
6     ...  
7 
```

Building

Granular control: shared memory OpenMP through Kokkos

```
1 $ cmake .. -DMODEL=lphi4 -DDEVICE=KOKKOS -DCUDA=OFF  
2 ...  
3 
```

TempLat

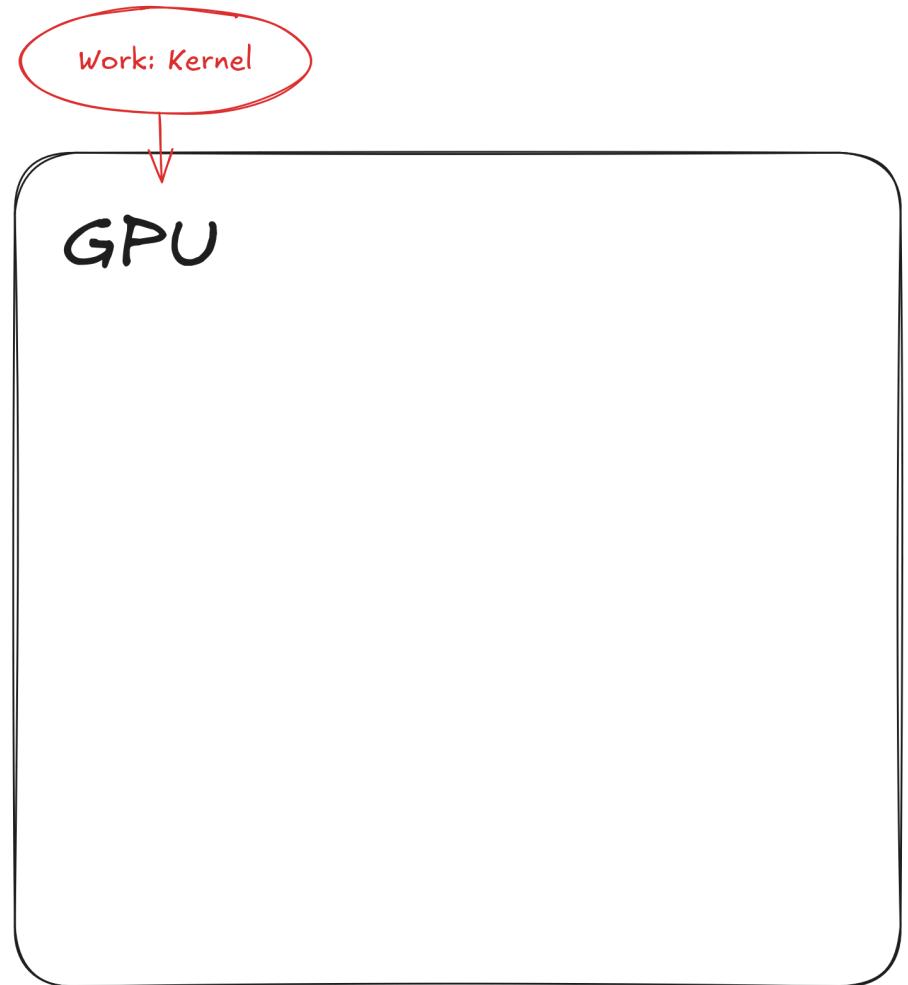
```
1 vType computeConfigurationSpace() {  
2     auto functor = DEVICE_CLASS_LAMBDA(const device::IdxArray<NDim> &idx,  
3                                         vType &update) {  
4         device::apply([&](const auto &...args) {  
5             update += GetValue::get(mT, args...);  
6         },  
7             idx);  
8     };  
9  
10    vType localResult{};  
11    device::iteration::reduce("Averager", cLayout, functor, localResult);  
12    return localResult;  
13 }
```

For "average" user:

Only minimal
changes

GPU architecture

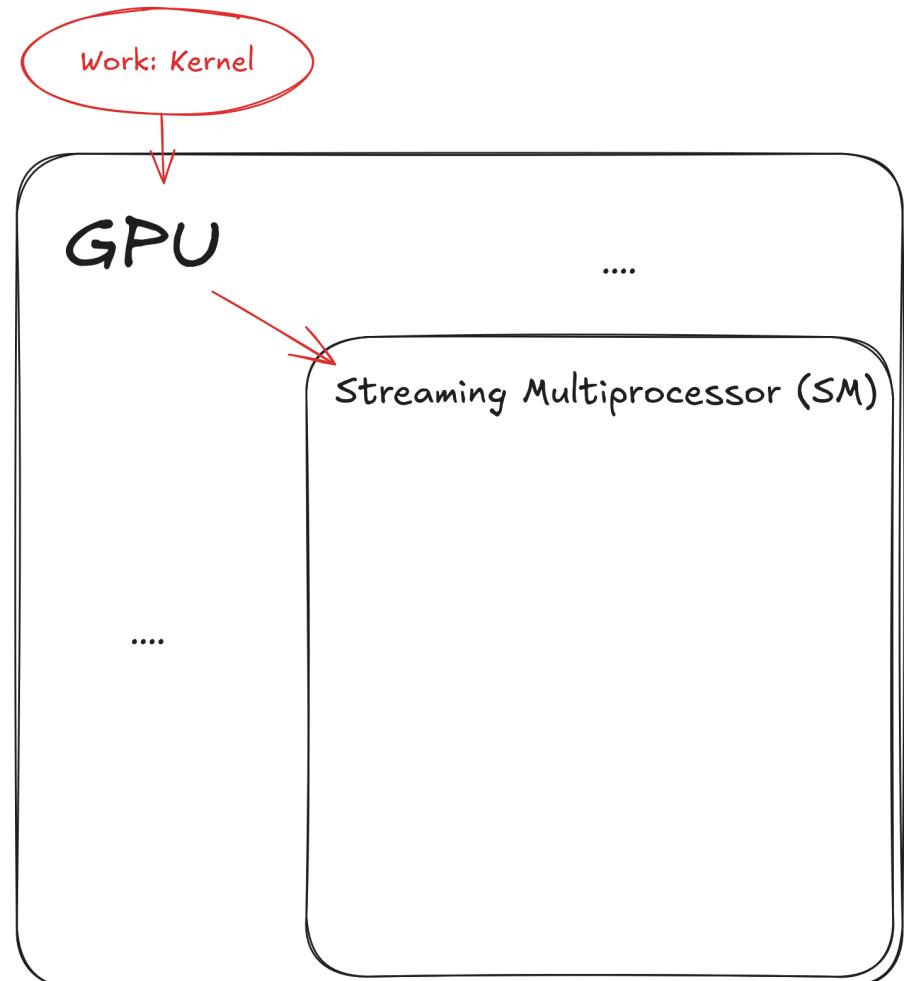
GPU thread hierarchy



GPU architecture

GPU thread hierarchy

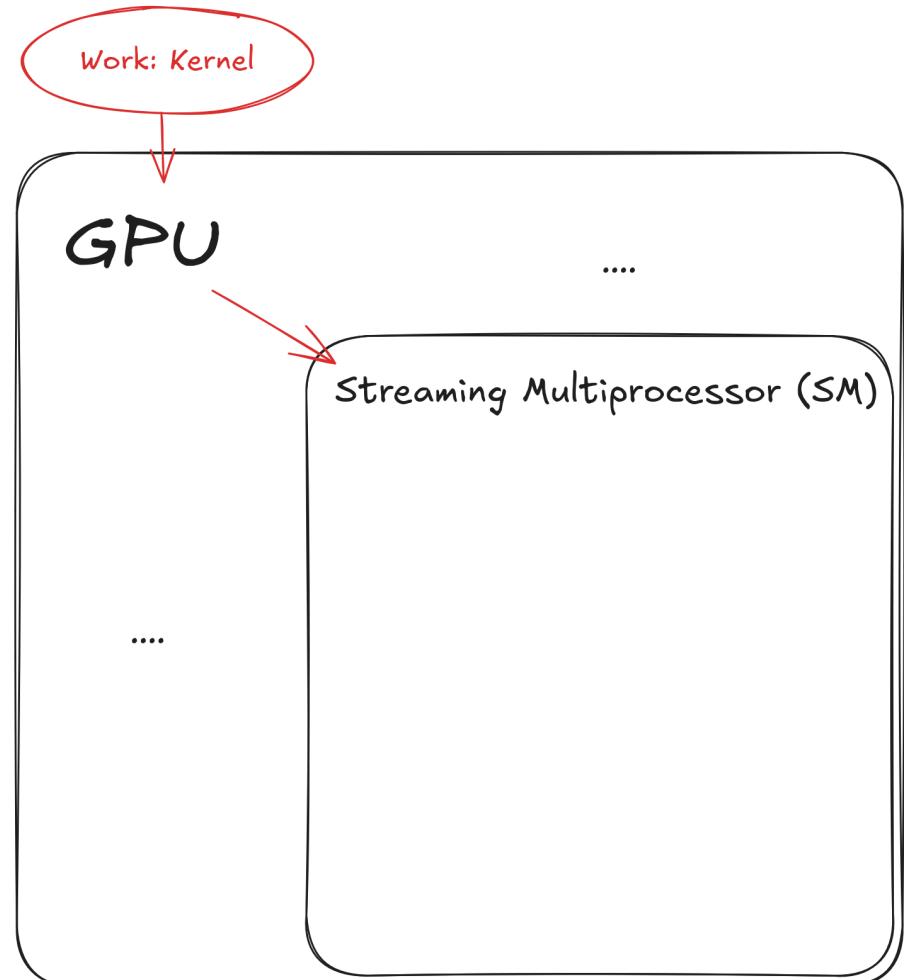
- GPUs have $\mathcal{O}(10)$ *Streaming multiprocessors*.



GPU architecture

GPU thread hierarchy

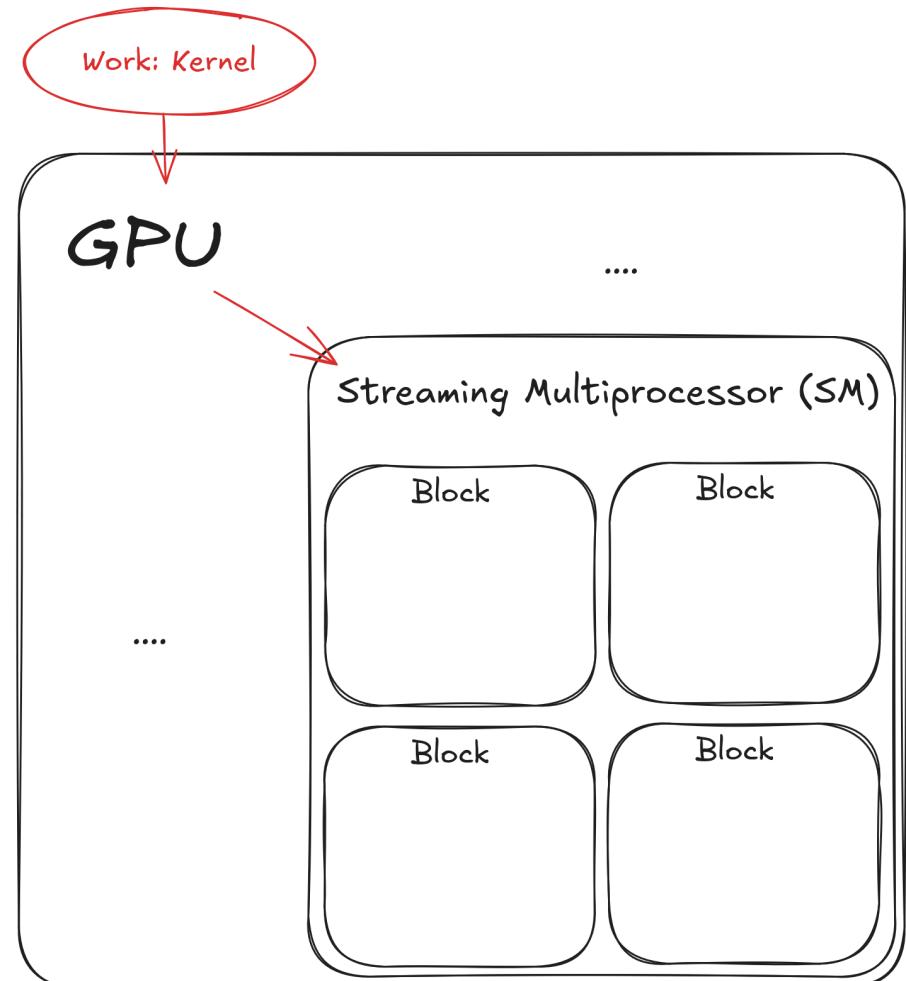
- GPUs have $\mathcal{O}(10)$ *Streaming multiprocessors*.
 - SMs execute *kernels* with series of parallel instructions.
 - SMs schedule their execution.



GPU architecture

GPU thread hierarchy

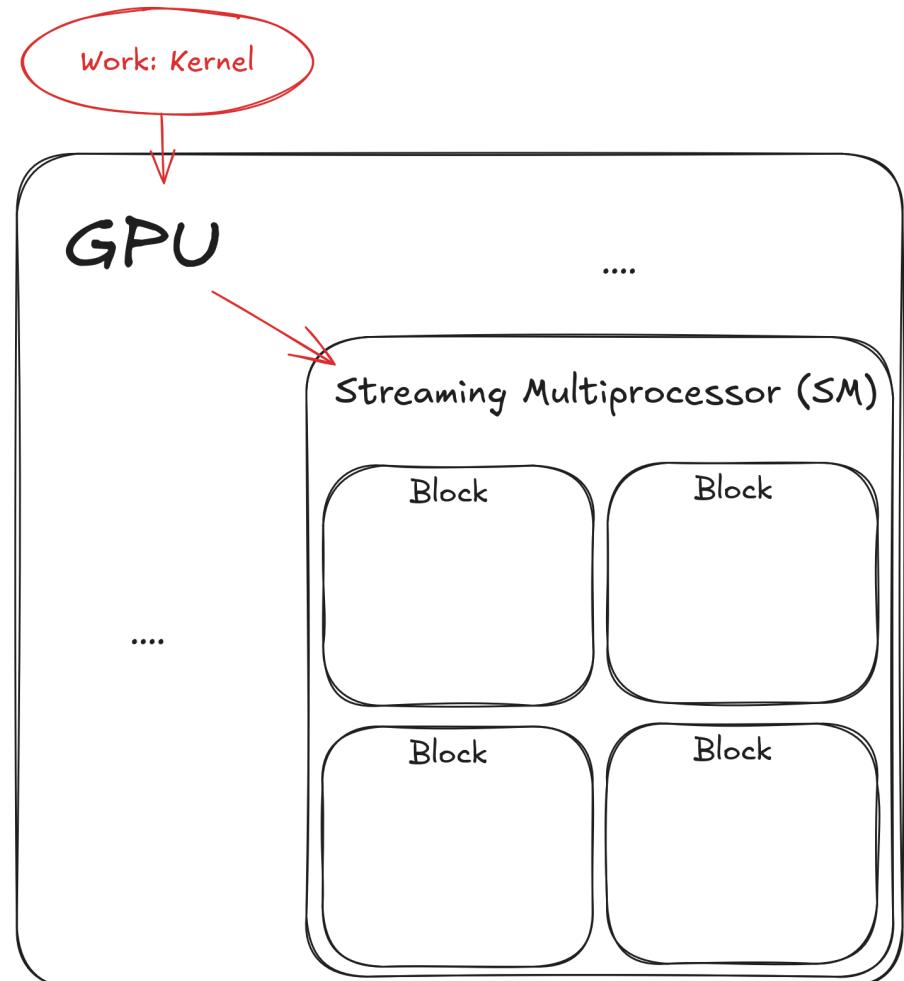
- GPUs have $\mathcal{O}(10)$ *Streaming multiprocessors*.
 - SMs execute *kernels* with series of parallel instructions.
 - SMs schedule their execution.
- Work given to SMs in *blocks*.



GPU architecture

GPU thread hierarchy

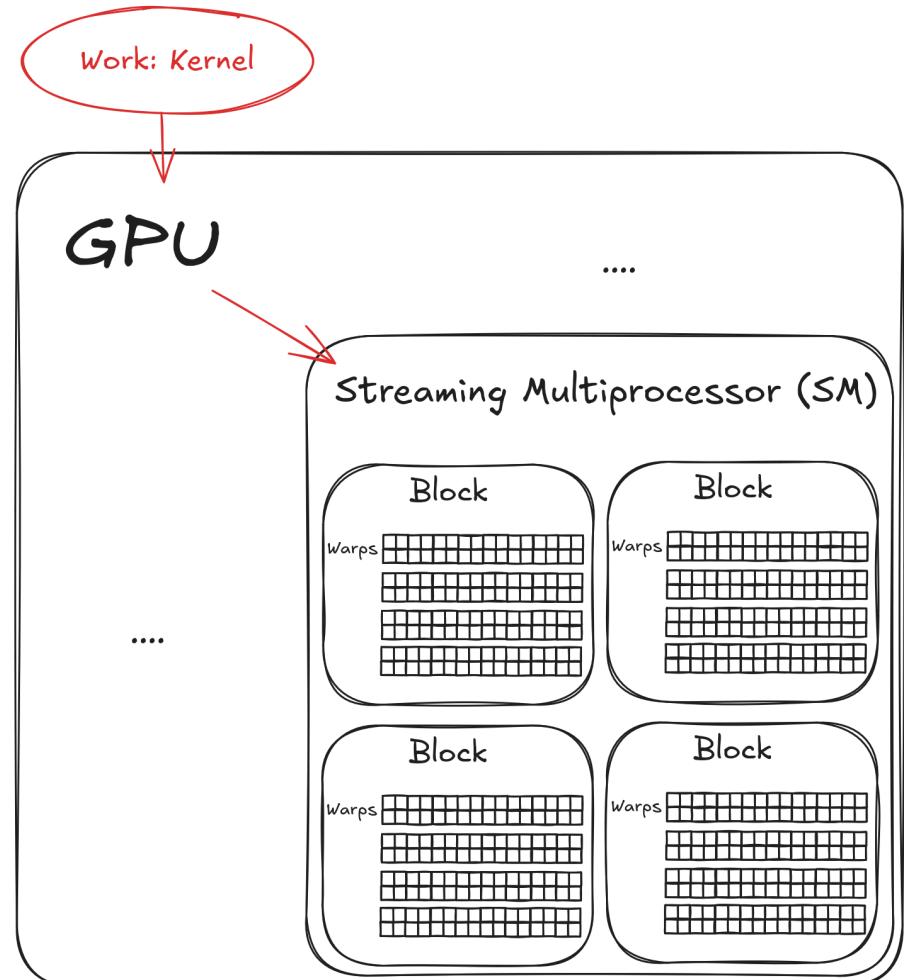
- GPUs have $\mathcal{O}(10)$ *Streaming multiprocessors*.
 - SMs execute *kernels* with series of parallel instructions.
 - SMs schedule their execution.
- Work given to SMs in *blocks*.
 - A block has to fit onto a SM's hardware capabilities (~1024 threads/block).
 - Each block's (sub-)contexts are persistent throughout its execution.



GPU architecture

GPU thread hierarchy

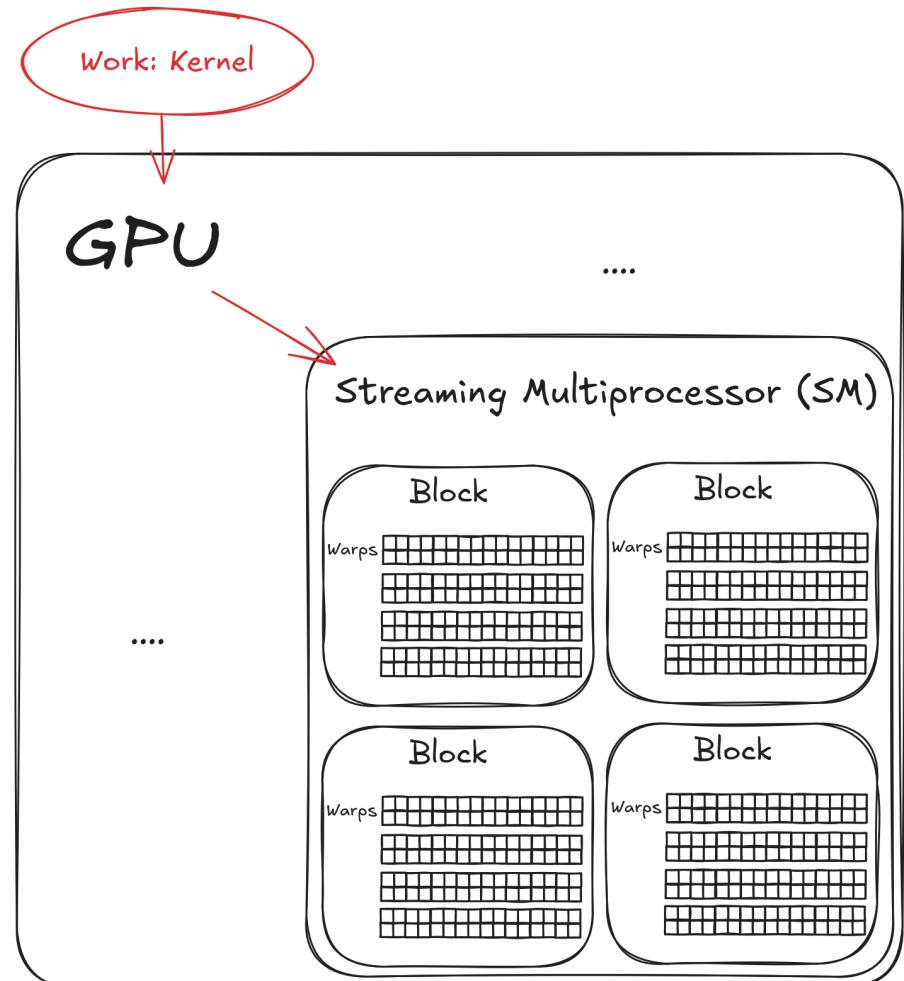
- GPUs have $\mathcal{O}(10)$ *Streaming multiprocessors*.
 - SMs execute *kernels* with series of parallel instructions.
 - SMs schedule their execution.
- Work given to SMs in *blocks*.
 - A block has to fit onto a SM's hardware capabilities (~1024 threads/block).
 - Each block's (sub-)contexts are persistent throughout its execution.
- Internally, blocks are subdivided into *warp*s.



GPU architecture

GPU thread hierarchy

- GPUs have $\mathcal{O}(10)$ *Streaming multiprocessors*.
 - SMs execute *kernels* with series of parallel instructions.
 - SMs schedule their execution.
- Work given to SMs in *blocks*.
 - A block has to fit onto a SM's hardware capabilities (~1024 threads/block).
 - Each block's (sub-)contexts are persistent throughout its execution.
- Internally, blocks are subdivided into *warps*.
 - Each warp runs a single instruction in a *kernel* in parallel.
 - Warp size is always 32 for Nvidia, 32 or 64 for AMD.



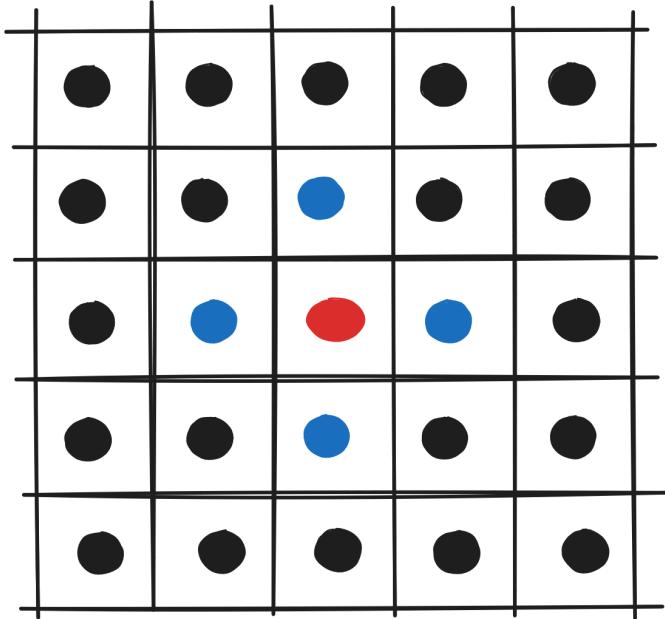
Memory access patterns

Coalescing vs. sequential access

Example: Solving massless Klein-Gordon equation in $d = 3$,

$$\partial_t^2 \phi(t, x) = \Delta \phi(t, x).$$

- Calculation of 1 thread at **red site**.
- **Blue sites** dependents for lattice Laplacian.



Memory access patterns

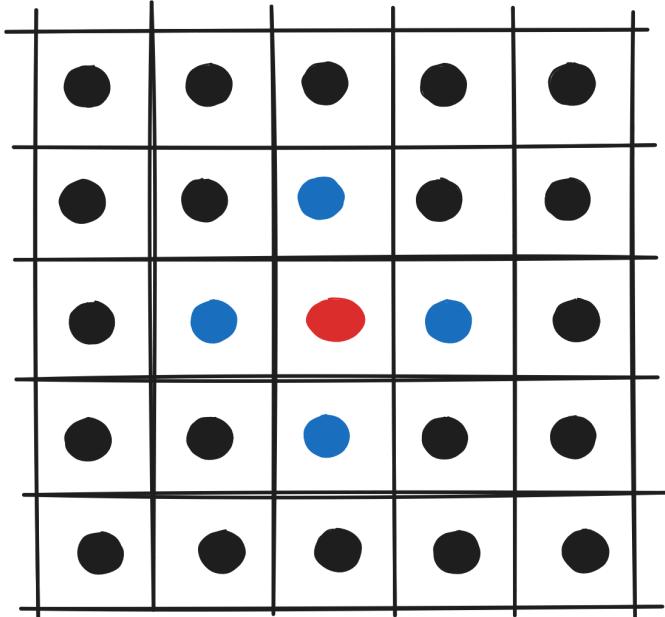
Coalescing vs. sequential access

Example: Solving massless Klein-Gordon equation in $d = 3$,

$$\partial_t^2 \phi(t, x) = \Delta \phi(t, x).$$

- Calculation of 1 thread at **red site**.
- **Blue sites** dependents for lattice Laplacian.

What is the optimal way to iterate over sites?



Memory access patterns

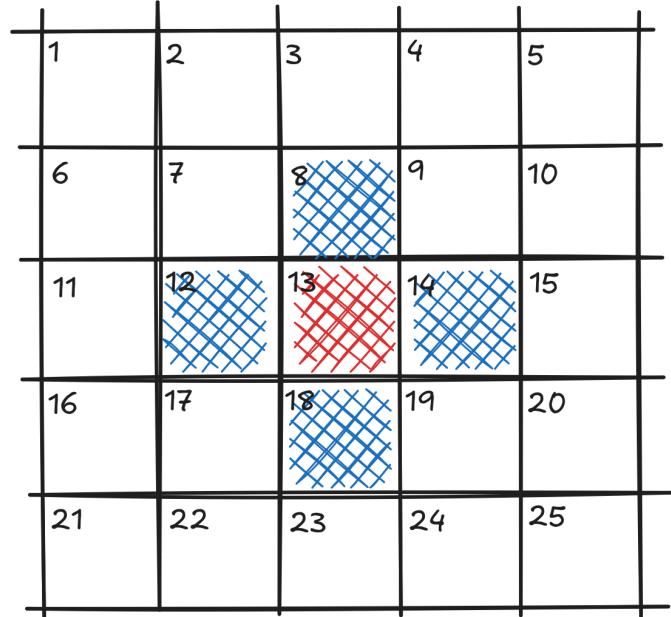
Coalescing vs. sequential access

Example: Solving massless Klein-Gordon equation in $d = 3$,

$$\partial_t^2 \phi(t, x) = \Delta \phi(t, x).$$

- Calculation of 1 thread at **red site**.
- **Blue sites** dependents for lattice Laplacian.
- Memory is ordered row-major.

What is the optimal way to iterate over sites?



Memory access patterns

Coalescing vs. sequential access

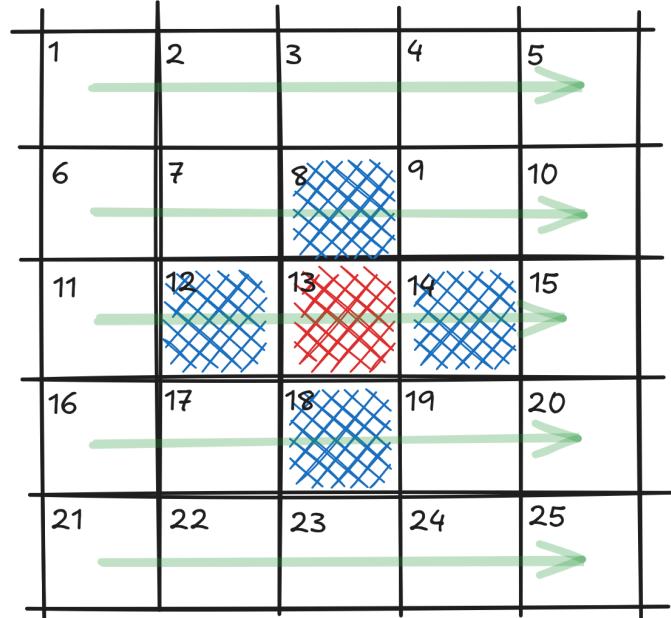
Example: Solving massless Klein-Gordon equation in $d = 3$,

$$\partial_t^2 \phi(t, x) = \Delta \phi(t, x).$$

- Calculation of 1 thread at **red site**.
- **Blue sites** dependents for lattice Laplacian.
- Memory is ordered row-major.

What is the optimal way to iterate over sites?

CPU: **Cashed access pattern** allows for caching of subsequent operations of a single thread.



Memory access patterns

Coalescing vs. sequential access

Example: Solving massless Klein-Gordon equation in $d = 3$,

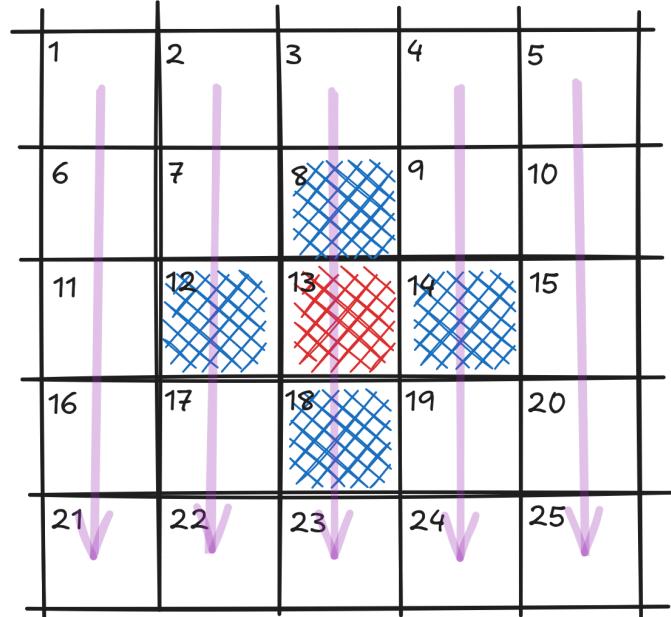
$$\partial_t^2 \phi(t, x) = \Delta \phi(t, x).$$

- Calculation of 1 thread at **red site**.
- **Blue sites** dependents for lattice Laplacian.
- Memory is ordered row-major.

What is the optimal way to iterate over sites?

CPU: **Cashed access pattern** allows for caching of subsequent operations of a single thread.

GPU: **Coalesced access pattern** allows for simultaneous reading of memory for multiple threads.



Memory access patterns

Coalescing vs. sequential access

Example: Solving massless Klein-Gordon equation in $d = 3$,

$$\partial_t^2 \phi(t, x) = \Delta \phi(t, x).$$

- Calculation of 1 thread at **red site**.
- **Blue sites** dependents for lattice Laplacian.
- Memory is ordered row-major.

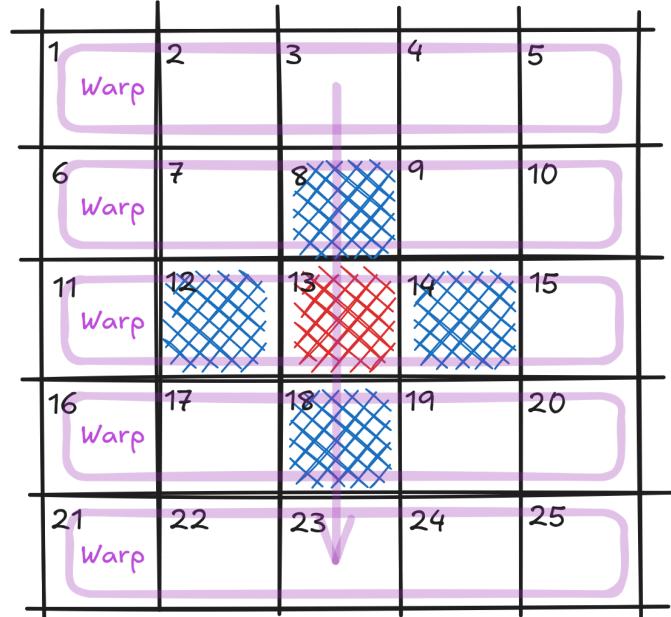
What is the optimal way to iterate over sites?

CPU: **Cashed access pattern** allows for caching of subsequent operations of a single thread.

GPU: **Coalesced access pattern** allows for simultaneous reading of memory for multiple threads.

This is similar to vectorization on a CPU!

SIMD (Single Instruction, Multiple Data) vs **SIMT** (Single Instruction, Multiple Threads)



Memory access patterns

Coalesced vs. cached access

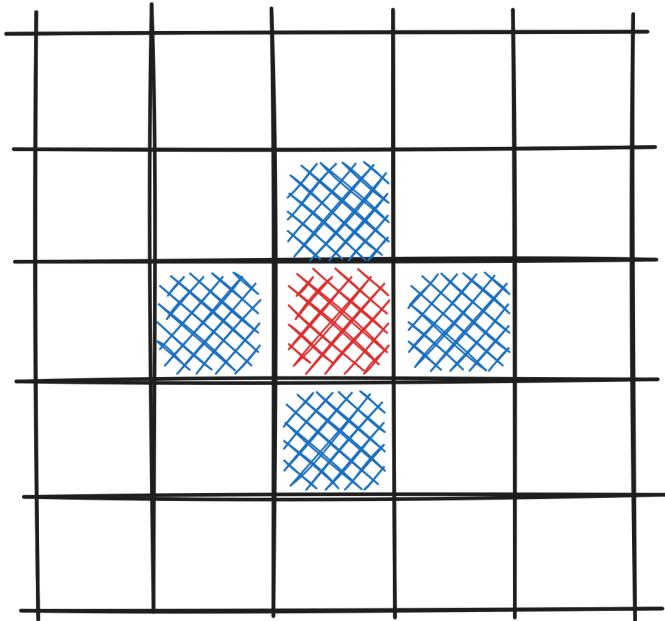
Example: Solving massless Klein-Gordon equation in $d = 3$,

$$\partial_t^2 \phi(t, x) = \Delta \phi(t, x).$$

- Calculation of 1 thread at **red site**.
- **Blue sites** dependents for lattice Laplacian.
- Memory is ordered row-major.

CPU: Prefer **prefer row-major access pattern**.

GPU: Prefer **column-major access pattern**.



How does this perform *in vivo*?

```
1 #define FORCE_ACCESS_PATTERN 0 // or 1
2 ...
3
4 int main(int argc, char **argv)
5 {
6     constexpr size_t NDim = 3;
7     using T = double;
8     constexpr size_t nGrid = 512;
9     constexpr size_t nGhost = 1;
10    constexpr size_t nSteps = 512;
11    constexpr T dt = 0.01;
12    ...
13    Field<NDim, T> phi("phi", toolBox);
14    Field<NDim, T> pi("pi", toolBox);
15
16    Benchmark bench([&](Benchmark::Measurer &measurer) {
17        phi.inFourierSpace() = RandomGaussianField<NDim, T>("Rand", toolBox);
18        pi.inFourierSpace() = RandomGaussianField<NDim, T>("Rand2", toolBox);
19
20        for (size_t i = 0; i < nSteps; ++i) {
21            pi.updateGhosts();
22            device::iteration::fence();
23            measurer.measure("timestepping", [&]() {
24                pi = pi + dt * LatticeLaplacian<NDim, decltype(phi)>(phi); // kick
25                phi = phi + dt * pi; // drift
26                device::iteration::fence();
27            });
28        });
29    });
30}
```

```
1 ...
2     for (size_t i = 0; i < nSteps; ++i) {
3         pi.updateGhosts();
4         device::iteration::fence();
5         measurer.measure("timestepping", [&](){
6             pi = pi + dt * LatticeLaplacian<NDim, decltype(phi)>(phi); // kick
7             phi = phi + dt * pi; // drift
8             device::iteration::fence();
9         });
10    }
```

Running this on my PC:

GPU: Nvidia 4070RTX mobile - 4788 Cores @ 2.175 GHz

CPU: Ryzen 9 7945HX - 16 Cores @ 5.4GHz

Taking a closer look

```
...
Benchmark bench([&](Benchmark::Measurer &measurer) {
    measurer.measure("x->k fourier", [&]() {
        phi.getMemoryManager()->confirmFourierSpace();
        pi.getMemoryManager()->confirmFourierSpace();
    });

    measurer.measure("initialize field", [&]() {
        phi.inFourierSpace() = RandomGaussianField<NDim, T>("Hoi", toolBox);
        pi.inFourierSpace() = RandomGaussianField<NDim, T>("Hai", toolBox);
    });

    measurer.measure("k->x fourier", [&]() {
        phi.getMemoryManager()->confirmConfigSpace();
        pi.getMemoryManager()->confirmConfigSpace();
    });

    for (size_t i = 0; i < nSteps; ++i) {
        measurer.measure("ghosts", [&]() {
            pi.updateGhosts();
            device::iteration::fence();
        });
        measurer.measure("timestepping", [&]() {
            pi = pi + dt * LatticeLaplacian<NDim, decltype(phi)>(phi);
        });
    }
});
```

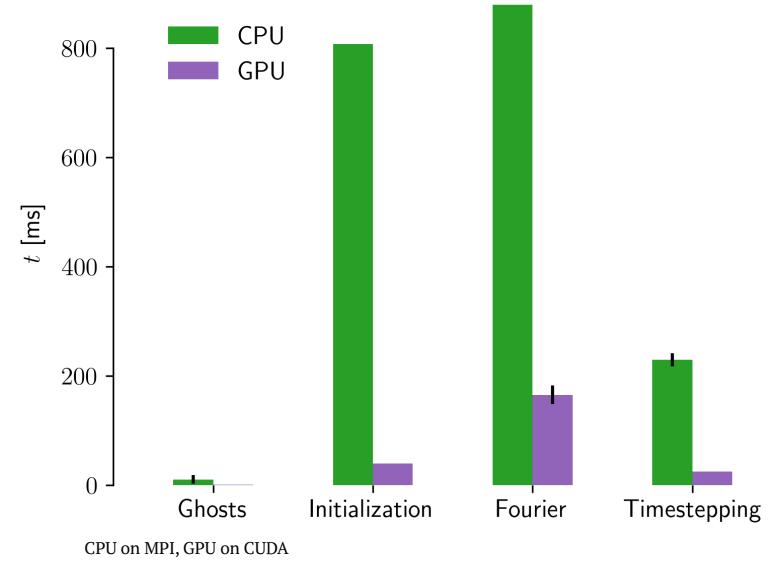
Taking a closer look

```
...
Benchmark bench([&](Benchmark::Measurer &measurer) {
    measurer.measure("x->k fourier", [&]() {
        phi.getMemoryManager()->confirmFourierSpace();
        pi.getMemoryManager()->confirmFourierSpace();
    });

    measurer.measure("initialize field", [&]() {
        phi.inFourierSpace() = RandomGaussianField<NDim, T>("Hoi", toolBox);
        pi.inFourierSpace() = RandomGaussianField<NDim, T>("Hai", toolBox);
    });

    measurer.measure("k->x fourier", [&]() {
        phi.getMemoryManager()->confirmConfigSpace();
        pi.getMemoryManager()->confirmConfigSpace();
    });

    for (size_t i = 0; i < nSteps; ++i) {
        measurer.measure("ghosts", [&]() {
            pi.updateGhosts();
            device::iteration::fence();
        });
        measurer.measure("timestepping", [&]() {
            pi = pi + dt * LatticeLaplacian<NDim, decltype(phi)>(phi);
        });
    }
});
```



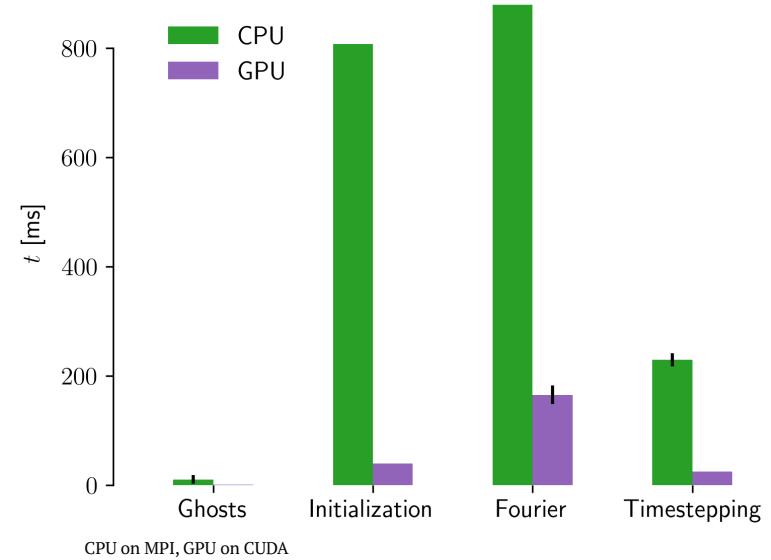
Taking a closer look

```
...
Benchmark bench([&](Benchmark::Measurer &measurer) {
    measurer.measure("x->k fourier", [&]() {
        phi.getMemoryManager()->confirmFourierSpace();
        pi.getMemoryManager()->confirmFourierSpace();
    });

    measurer.measure("initialize field", [&]() {
        phi.inFourierSpace() = RandomGaussianField<NDim, T>("Hoi", toolBox);
        pi.inFourierSpace() = RandomGaussianField<NDim, T>("Hai", toolBox);
    });

    measurer.measure("k->x fourier", [&]() {
        phi.getMemoryManager()->confirmConfigSpace();
        pi.getMemoryManager()->confirmConfigSpace();
    });

    for (size_t i = 0; i < nSteps; ++i) {
        measurer.measure("ghosts", [&]() {
            pi.updateGhosts();
            device::iteration::fence();
        });
        measurer.measure("timestepping", [&]() {
            pi = pi + dt * LatticeLaplacian<NDim, decltype(phi)>(phi);
        });
    }
});
```



- Fourier transformation: cuFFT
(automatic switch to GPU native FFTs)

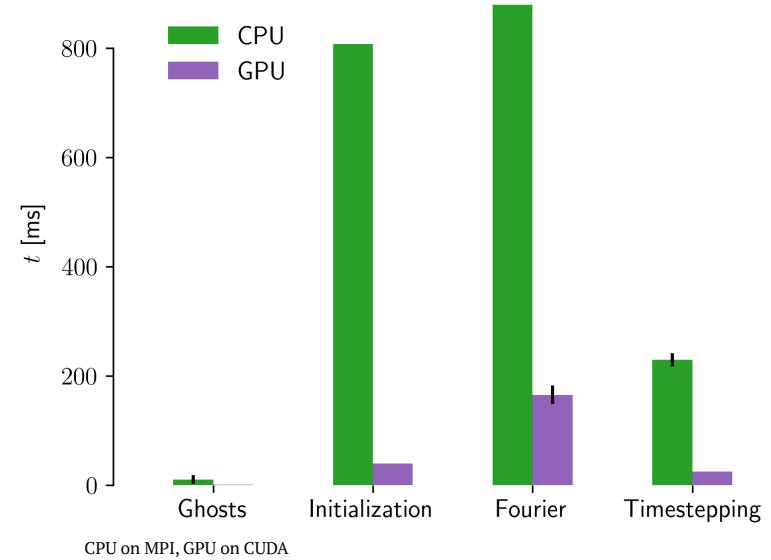
Taking a closer look

```
...
Benchmark bench([&](Benchmark::Measurer &measurer) {
    measurer.measure("x->k fourier", [&]() {
        phi.getMemoryManager()->confirmFourierSpace();
        pi.getMemoryManager()->confirmFourierSpace();
    });

    measurer.measure("initialize field", [&]() {
        phi.inFourierSpace() = RandomGaussianField<NDim, T>("Hoi", toolBox);
        pi.inFourierSpace() = RandomGaussianField<NDim, T>("Hai", toolBox);
    });

    measurer.measure("k->x fourier", [&]() {
        phi.getMemoryManager()->confirmConfigSpace();
        pi.getMemoryManager()->confirmConfigSpace();
    });

    for (size_t i = 0; i < nSteps; ++i) {
        measurer.measure("ghosts", [&]() {
            pi.updateGhosts();
            device::iteration::fence();
        });
        measurer.measure("timestepping", [&]() {
            pi = pi + dt * LatticeLaplacian<NDim, decltype(phi)>(phi);
        });
    }
});
```



- Fourier transformation: cuFFT
(automatic switch to GPU native FFTs)

Benchmarking $\lambda\phi^4$ -theory

with the `lphi4` model in CosmoLattice

Benchmarking Λ -theory

with the lphi4 model in CosmoLattice

PRELIMINARY

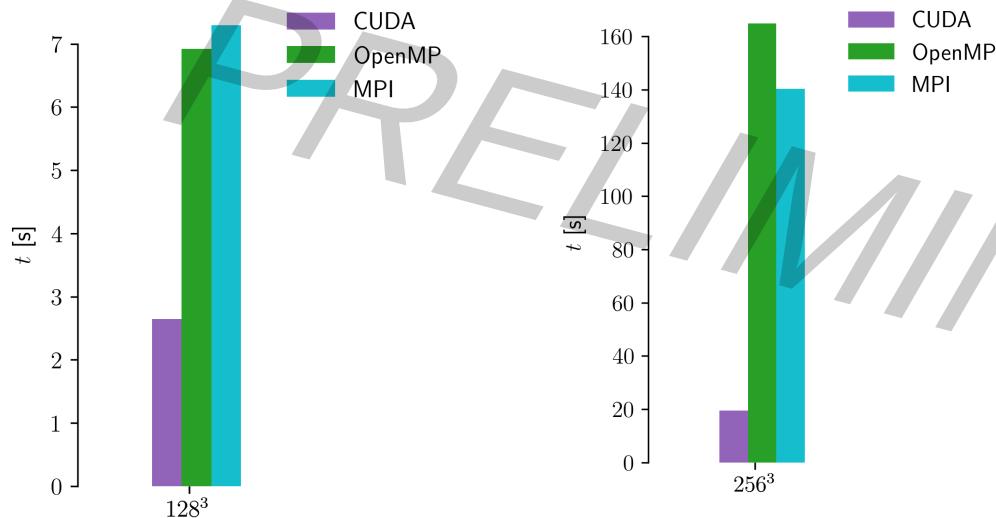
Benchmarking -theory

with the `lphi4` model in CosmoLattice



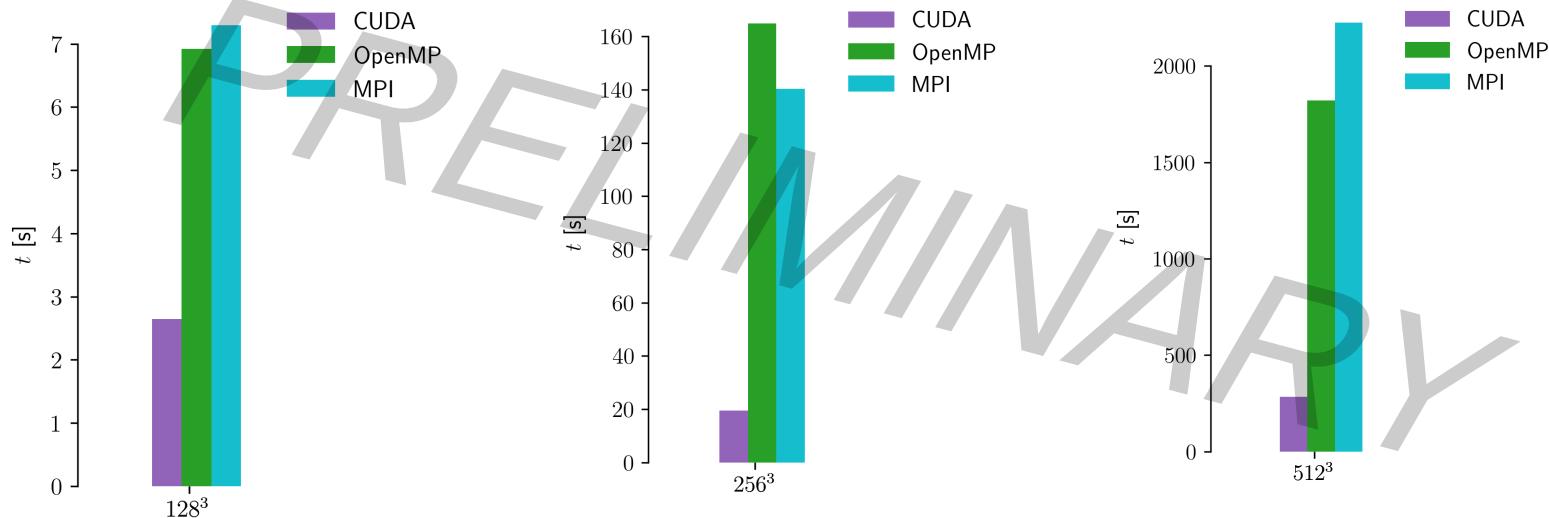
Benchmarking -theory

with the `lphi4` model in CosmoLattice



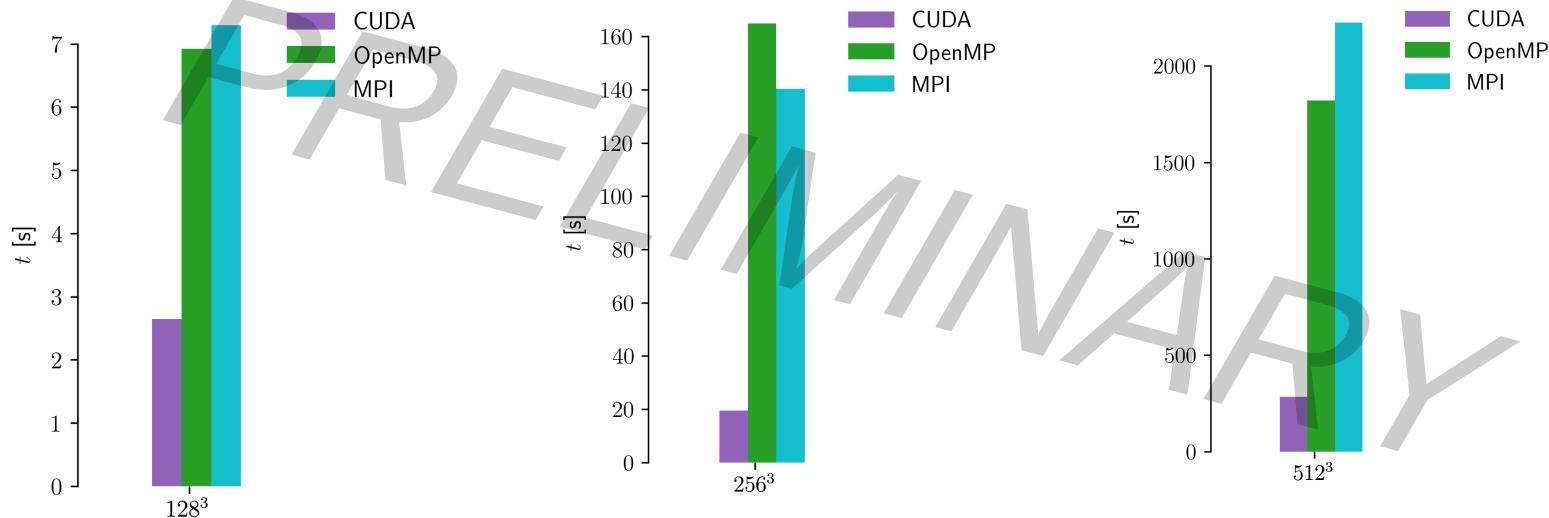
Benchmarking -theory

with the `lphi4` model in CosmoLattice



Benchmarking -theory

with the `lphi4` model in CosmoLattice



Slightly unfair comparison
(my CPU is "stronger")

Benchmarking -theory

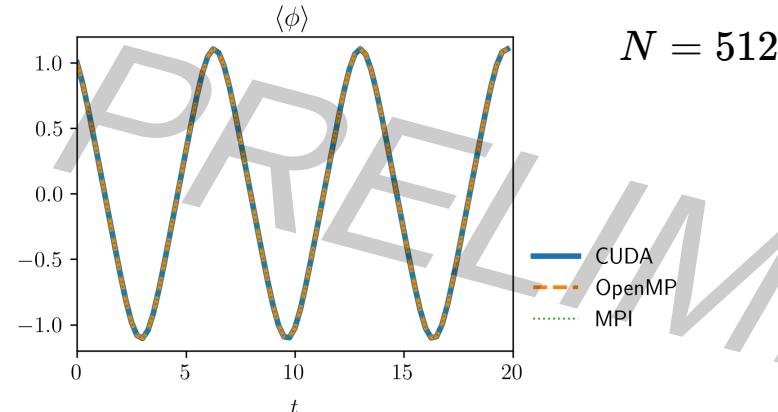
with the lphi4 model in CosmoLattice

$N = 512$

PRELIMINARY

Benchmarking ϕ -theory

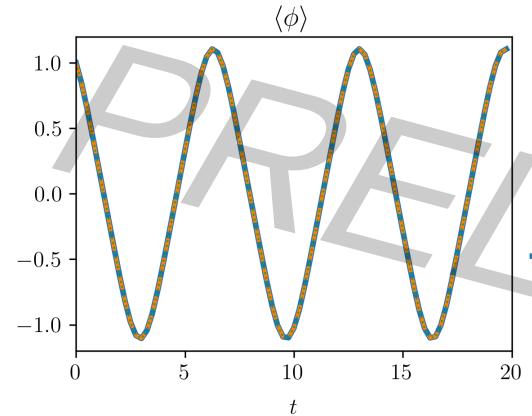
with the `lphi4` model in CosmoLattice



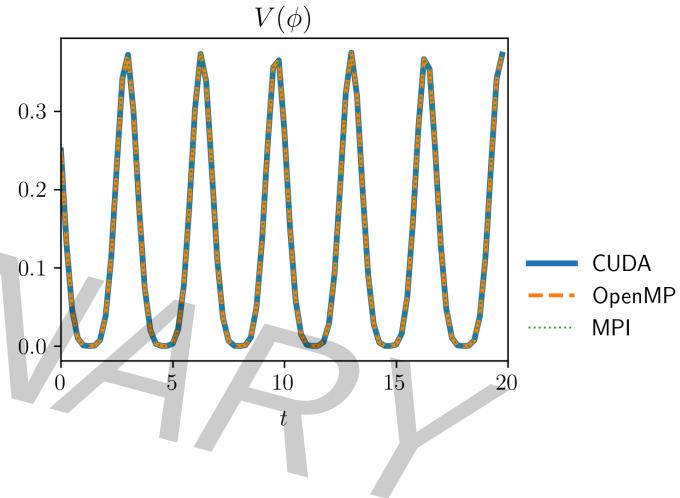
PRELIMINARY

Benchmarking ϕ -theory

with the `lphi4` model in CosmoLattice

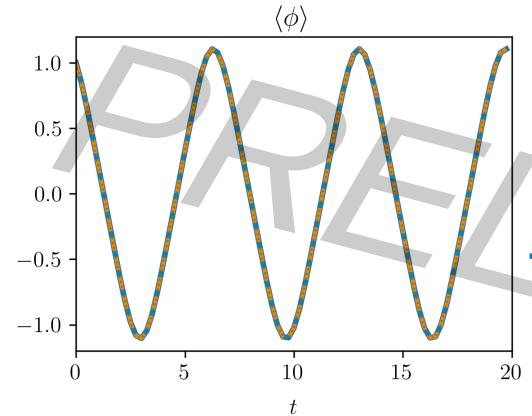


$N = 512$

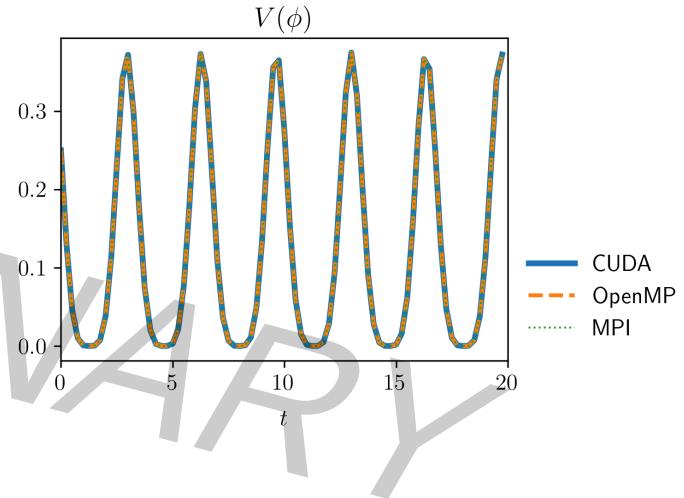


Benchmarking ϕ -theory

with the `lphi4` model in CosmoLattice



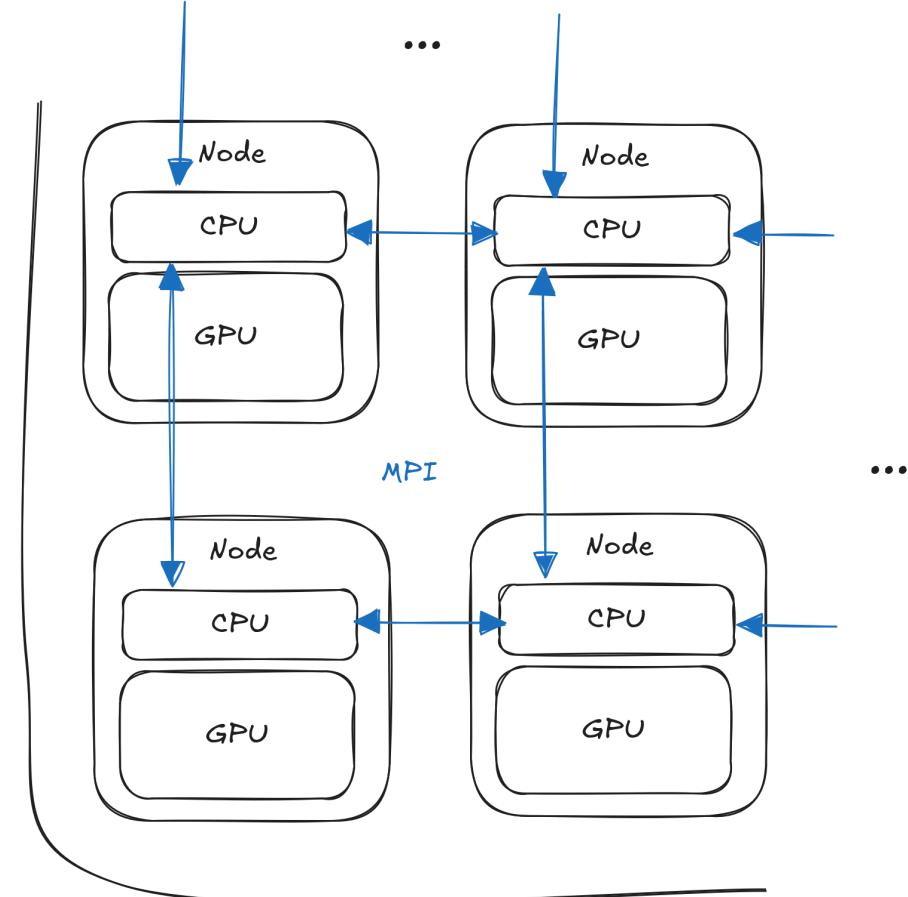
$N = 512$



Scaling it up

Using large GPU clusters

- To use large clusters and link up many nodes, CosmoLattice uses the **Message-Passing Interface (MPI)** (see lecture yesterday).
- Send data in RAM (e.g. ghosts) between neighbouring nodes.

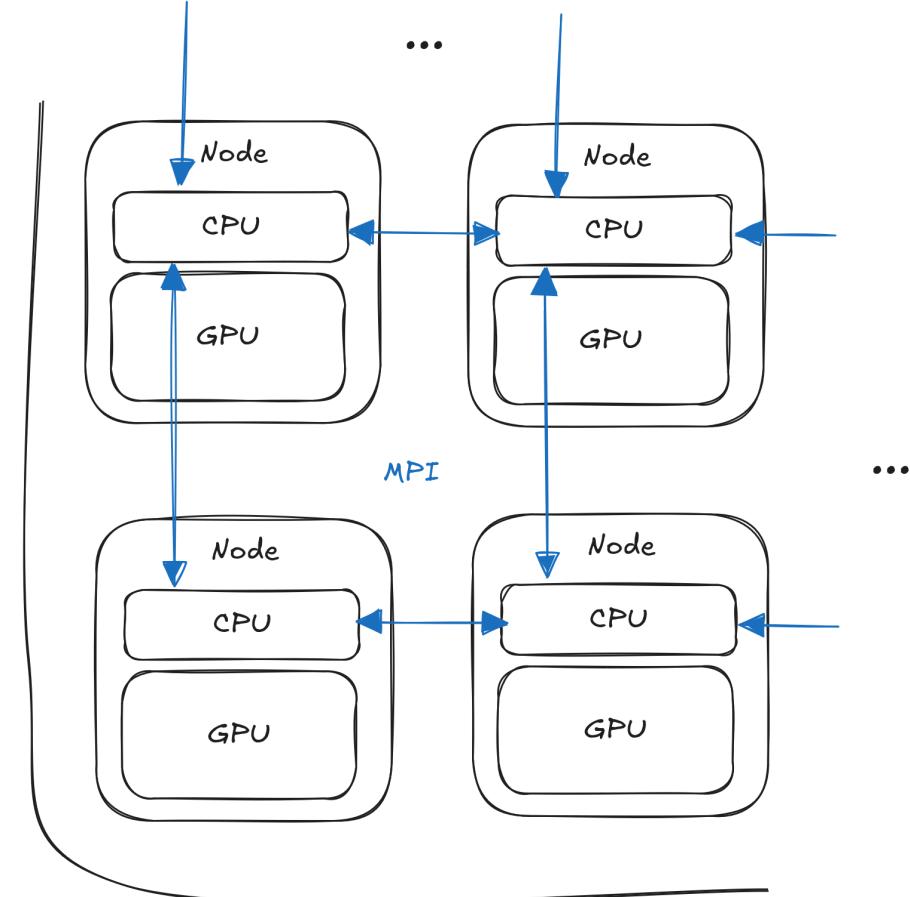


Scaling it up

Using large GPU clusters

- To use large clusters and link up many nodes, CosmoLattice uses the **Message-Passing Interface (MPI)** (see lecture yesterday).
- Send data in RAM (e.g. ghosts) between neighbouring nodes.

What about MPI+GPUs?



Scaling it up

Using large GPU clusters

- To use large clusters and link up many nodes, CosmoLattice uses the **Message-Passing Interface (MPI)** (see lecture yesterday).
- Send data in RAM (e.g. ghosts) between neighbouring nodes.

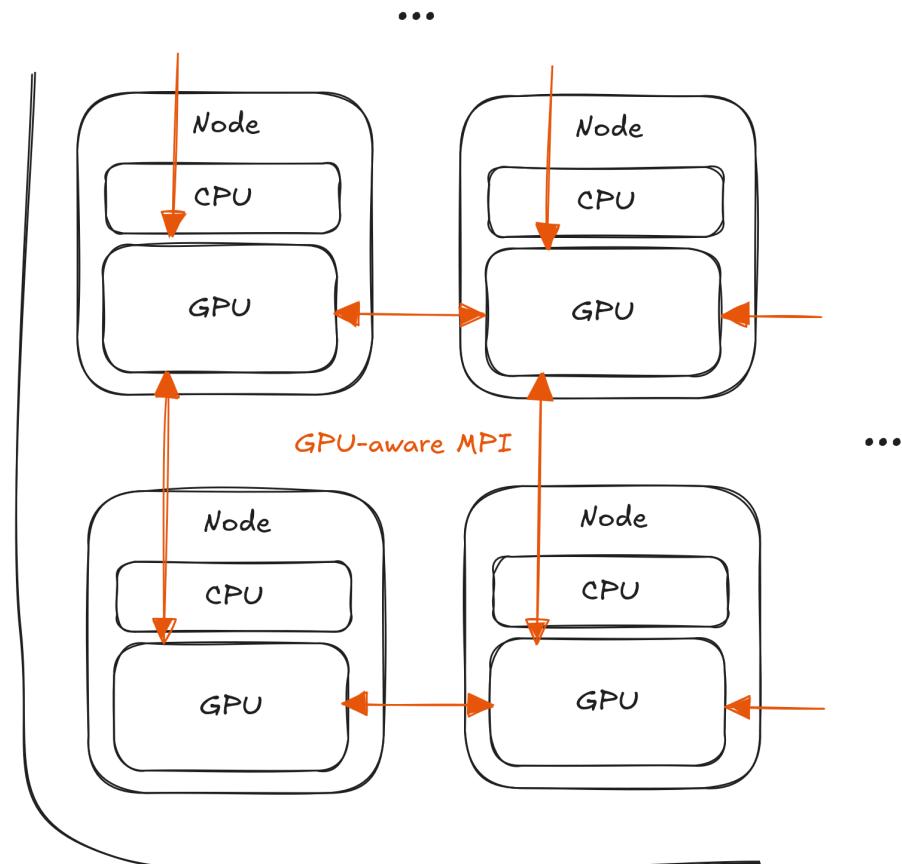
What about MPI+GPUs?

- GPU-aware MPI** can exchange data directly between device memory.

Support since before 2013:

- OpenMPI
- MVAPICH2
- Cray MPI
- IBM MPI

- No changes in MPI-code! except in FFT code...



There's more...

Useful, more fine-grained parallelization for even more speedup in the future:

There's more...

Useful, more fine-grained parallelization for even more speedup in the future:

- Hierarchical parallelism.
split complex tasks into smaller work-teams with the potential for caching

There's more...

Useful, more fine-grained parallelization for even more speedup in the future:

- Hierarchical parallelism.
split complex tasks into smaller work-teams with the potential for caching
- Block- and Warp- local caching.
reduce the number of repeated operations

There's more...

Useful, more fine-grained parallelization for even more speedup in the future:

- Hierarchical parallelism.
split complex tasks into smaller work-teams with the potential for caching
- Block- and Warp- local caching.
reduce the number of repeated operations
- Heterogeneous work-balancing.
e.g. use CPU for measurements while the GPU continues evolving

There's more...

Useful, more fine-grained parallelization for even more speedup in the future:

- Hierarchical parallelism.
split complex tasks into smaller work-teams with the potential for caching
- Block- and Warp- local caching.
reduce the number of repeated operations
- Heterogeneous work-balancing.
e.g. use CPU for measurements while the GPU continues evolving
- ...

Questions?

Thanks for your attention!

Release of CosmoLattice with GPUs ~ early 2026